

Probabilistic Information Retrieval Models

- Lecture 4
- SI 650 / EECS 549
- Information Retrieval
- Sept. 17, 2025

Some slides based on slide deck by Dr. Jurgens

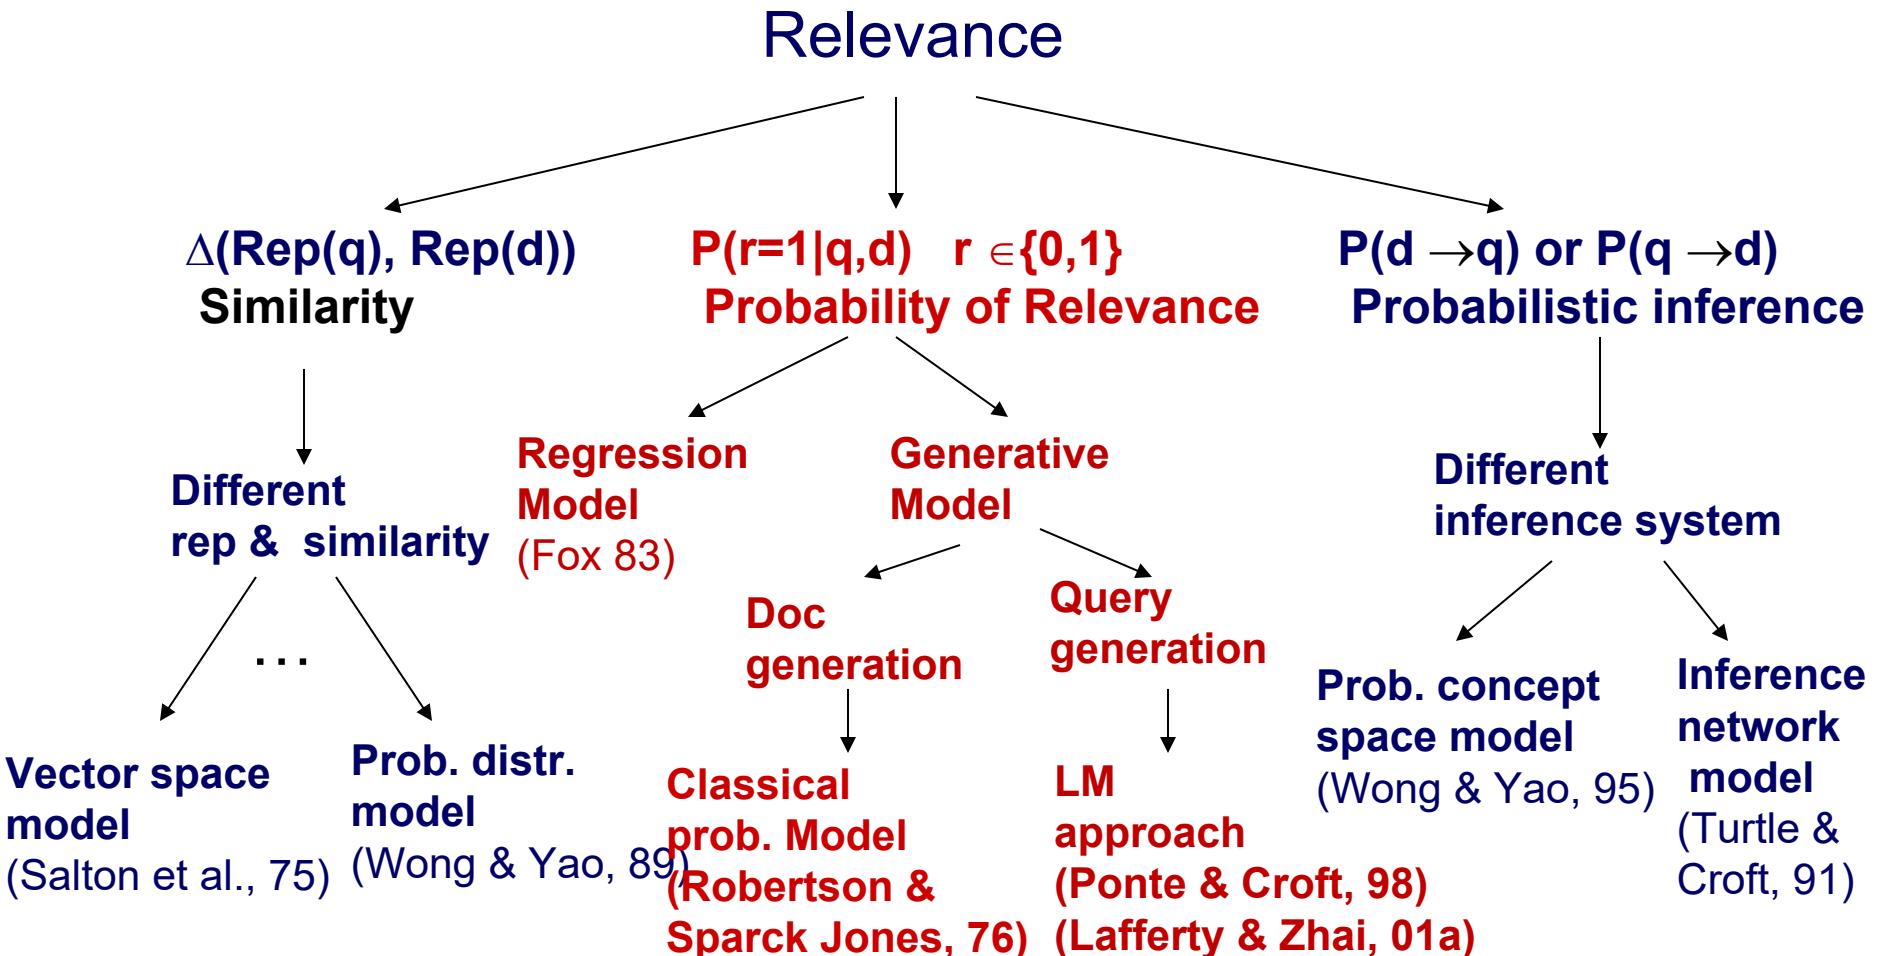
Administrative Notes

- Homework 1 due in two weeks!
 - Please try to go end-to-end first *before* implementing all of one part—get the IR system running!
- Project proposals due in three weeks! Attend discussion sections for more insights on this.

Today's Goals

- Think about retrieval beyond TF-IDF and VS Models
- Learn a bit more about probabilities

The Notion of Relevance



Prob/Statistics & Text Management

- Probability & statistics provide a principled way to quantify the uncertainties associated with natural language
- Allow us to answer questions like:
 - Given that we observe “baseball” three times and “game” once in a news article, how likely is it about “sports”?
(text categorization, information retrieval)
 - Given that a user is interested in sports news, how likely would the user use “baseball” in a query?
(information retrieval)

Things we will use in the course

- Events
- Probability, joint probability, conditional probability
- Distributions
- Bayes Rule
- Basic Statistics
- Parameter estimation
- Maximum likelihood estimation

Revisit: Essential Probability & Statistics

Random variable

- A variable that can take values within a fixed set (discrete) or within some range (continuous).

$$X \in \{1, 2, 3, 4, 5, 6\}$$

$$X \in \{\text{the}, \text{a}, \text{dog}, \text{cat}, \text{runs}, \text{to}, \text{store}\}$$

Example: Flip a Coin



- Sample space: {Head, Tail}
- Fair coin:
 - $p(H) = 0.5, p(T) = 0.5$
- Unfair coin, e.g.:
 - $p(H) = 0.3, p(T) = 0.7$
- Flipping two fair coins:
 - Sample space: {HH, HT, TH, TT}
- Modeling text:
 - Flip a coin to decide whether or not to include a word in a document
 - Sample space = {appear, absence}

Example: Toss a Dice



- Sample space: $S = \{1, 2, 3, 4, 5, 6\}$
- Fair dice:
 - $p(1) = p(2) = p(3) = p(4) = p(5) = p(6) = 1/6$
- Unfair dice: $p(1) = 0.3, p(2) = 0.2, \dots$
- N-dimensional dice:
 - $S = \{1, 2, 3, 4, \dots, N\}$
- Modeling text:
 - Toss a die to decide which word to write in the next position
 - Sample space = {cat, dog, tiger, ...}

What if the Dice has More Faces?

- Suitable to represent documents
- Every face corresponds to a word in vocabulary
- The author toss a dice to write a word
- Apparently, an unfair dice



Maximum Likelihood Estimate

- Data: a document d with counts $c(w_1), \dots, c(w_N)$, and length $|d|$
- Model: multinomial distribution M with parameters $\{p(w_i)\}$
$$p(w_i) = \frac{c(w_i)}{|d|}$$
- Likelihood: $p(d|M)$
- Maximum likelihood estimator:
 $M = \operatorname{argmax}_M p(d|M)$

Problem: a document is short
if $c(w_i) = 0$, does this mean $p(w_i) = 0$?

Revisit: Basic Concept of Probability

- Joint probability: $P(AB)$, also written as $P(A, B)$
- Conditional Probability: $P(B|A) = P(A, B)/P(A)$
 - $P(A, B) = P(A)P(B|A) = P(B)P(A|B)$
 - For **independent events**, $P(A, B) = P(A)P(B)$, and $P(A|B) = P(A)$, $P(B|A) = P(B)$.
- Bayes' Rule:
 - $P(A|B) = P(B|A)P(A)/P(B)$
 - $P(H|E) = P(E|H)P(H)/P(E)$
 - $P(C|D) = P(D|C)P(C)/P(D)$

H: hypothesis; E: Evidence

In text classification: C: class label; D: data (document)

Getting to Statistics ...

- We are flipping an unfair coin, but $P(\text{Head})=?$ (parameter estimation)
 - If we see the results of a huge number of random experiments, then
$$\hat{P}(\text{Head}) = \frac{\text{count(Heads)}}{\text{count(Flips)}} = \frac{\text{count(Heads)}}{\text{count(Heads)} + \text{count(Tails)}}$$
 - But, what if we only see a small sample (e.g., 2)? Is this estimate still reliable? We flip twice and got two tails, does it mean $P(\text{Head}) = 0?$
- In general, statistics has to do with drawing conclusions on the whole population based on **observations** of a sample (data)

Parameter Estimation of a Random Variable

- General setting:
 - Given a (hypothesized & probabilistic) model that governs the random experiment
 - The model gives a probability of any **data** $p(D|\theta)$ that depends on the **parameter** θ
 - Now, given actual sample data $X=\{x_1, \dots, x_n\}$, what can we say about the value of θ ?
- Intuitively, take your best guess of θ – “best” means “best explaining/fitting the data”
- Generally an optimization problem

$$P(X = x)$$

Probability that the random variable X takes the value x (e.g., 1)

$$X \in \{1, 2, 3, 4, 5, 6\}$$

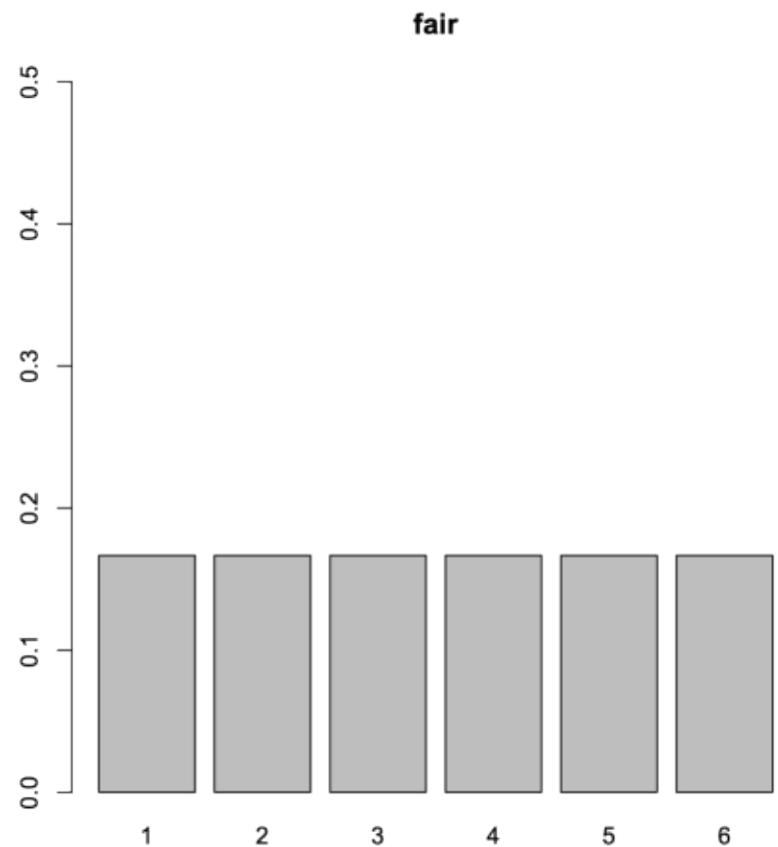
Two conditions:

1. Between 0 and 1: $0 \leq P(X = x) \leq 1$

2. Sum of all probabilities = 1 $\sum_x P(X = x) = 1$

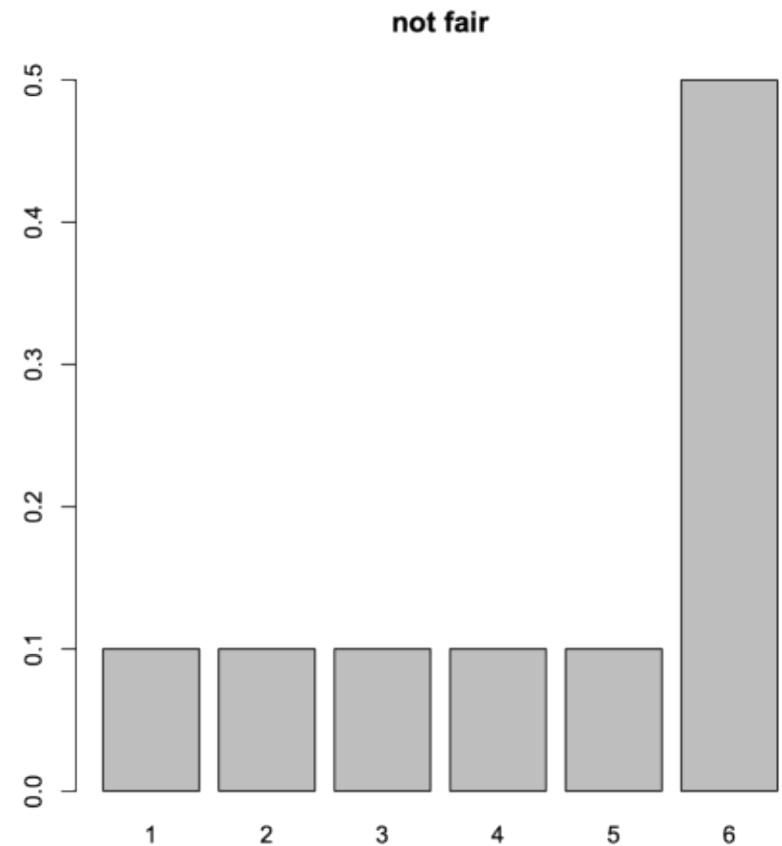
Fair dice

$$X \in \{1, 2, 3, 4, 5, 6\}$$



Weighted dice

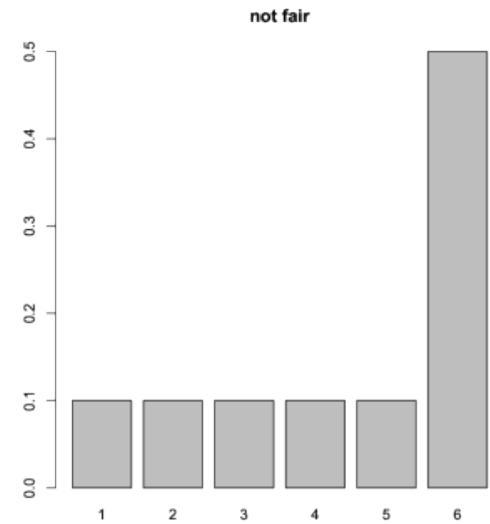
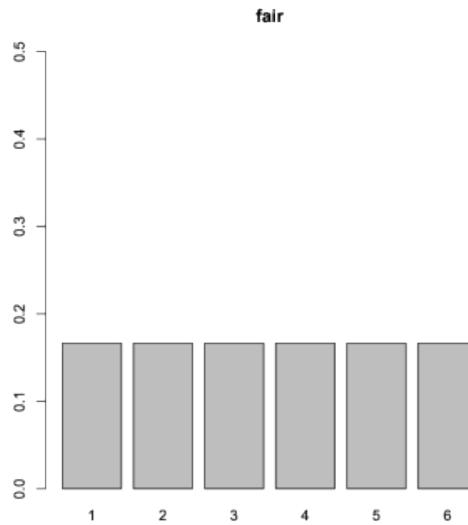
$$X \in \{1, 2, 3, 4, 5, 6\}$$



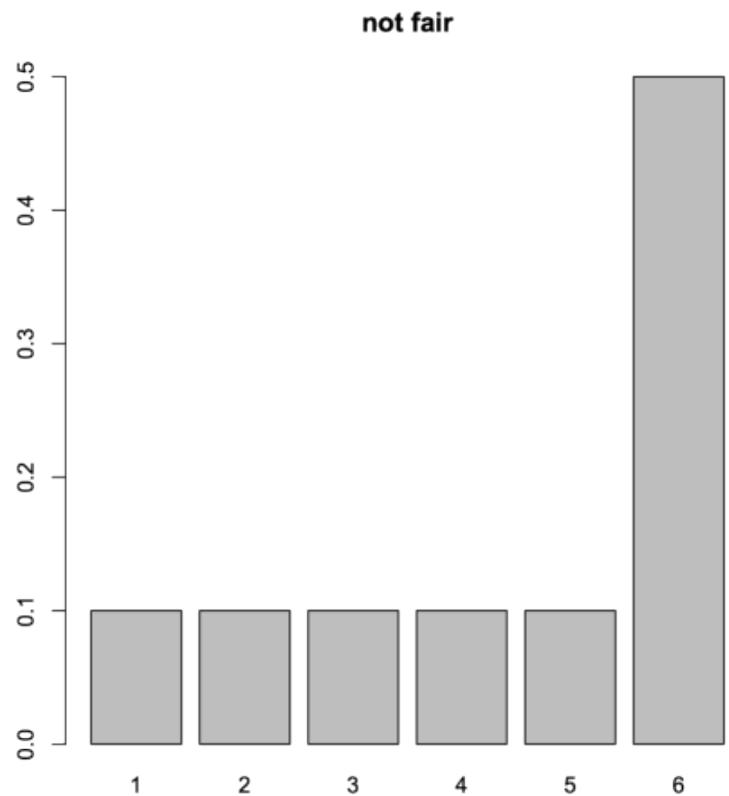
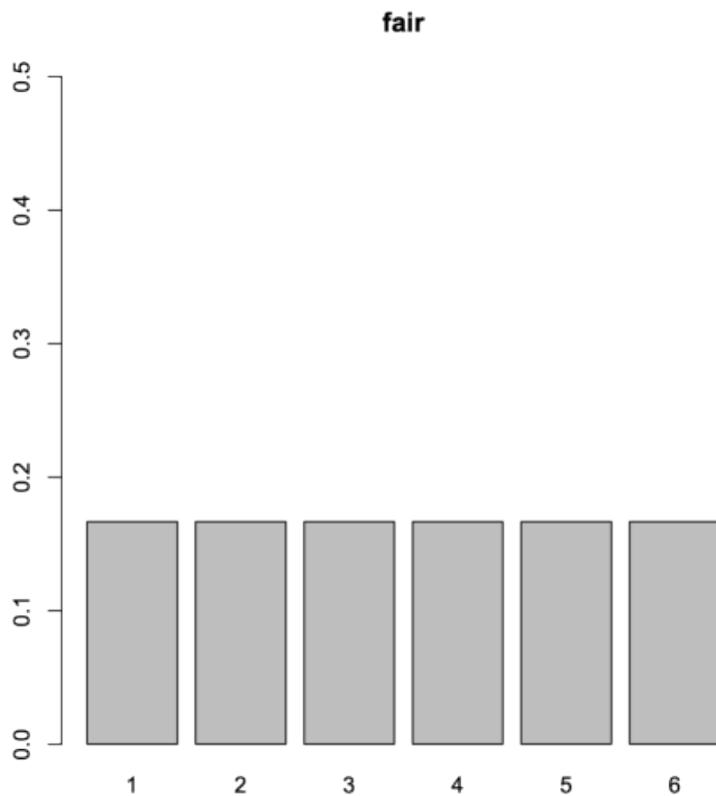
Inference

$$X \in \{1, 2, 3, 4, 5, 6\}$$

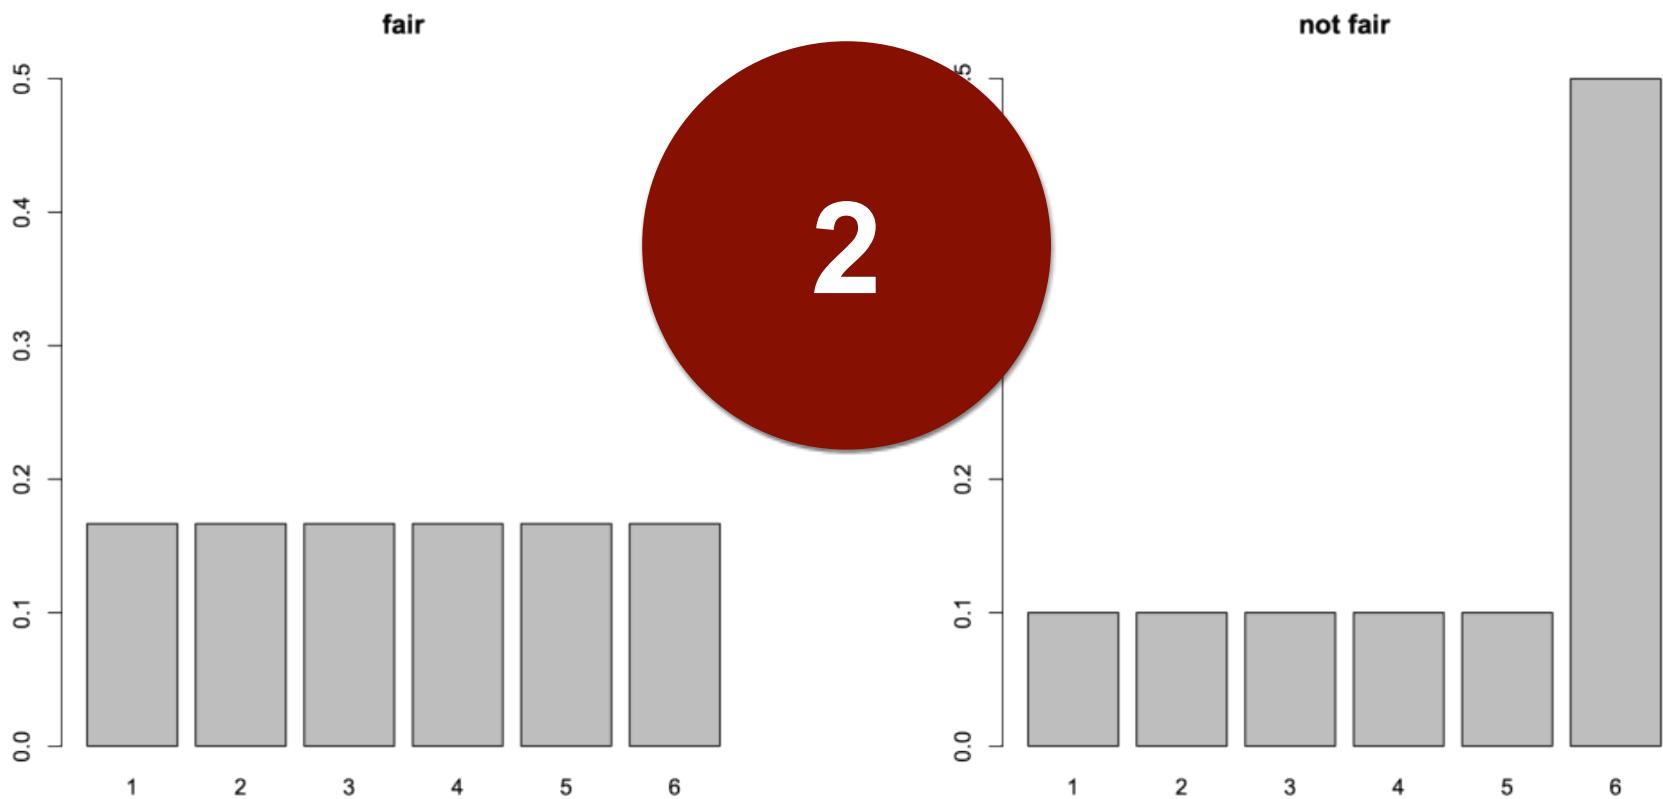
We want to *infer* the probability distribution that generated the data we see.



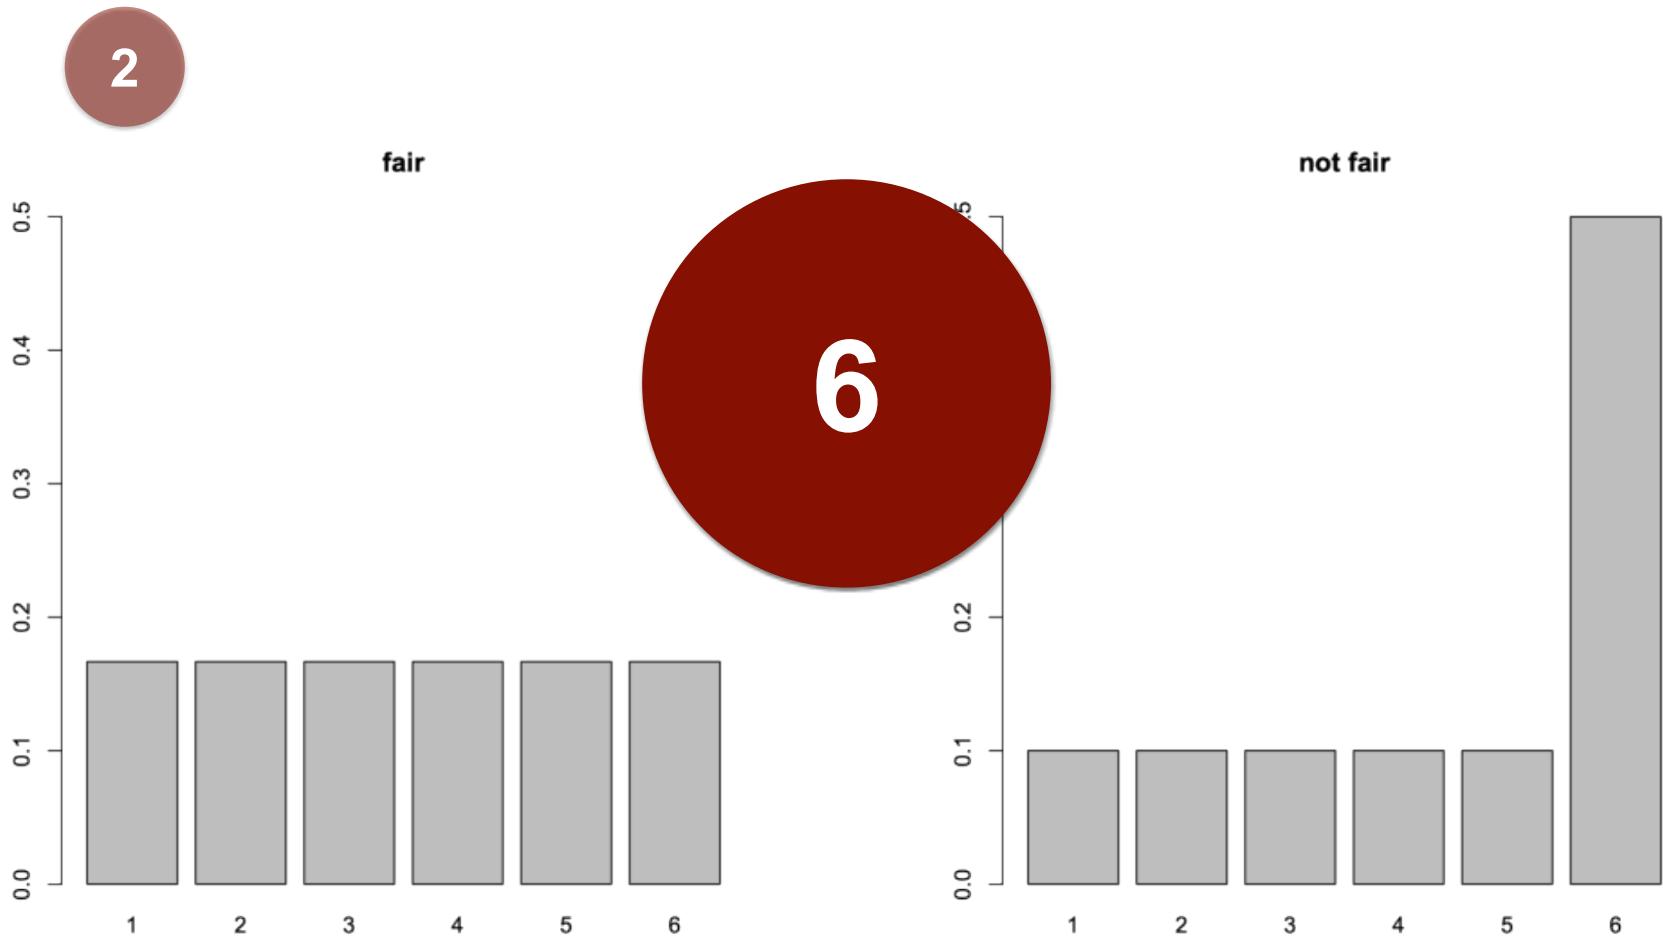
Probability



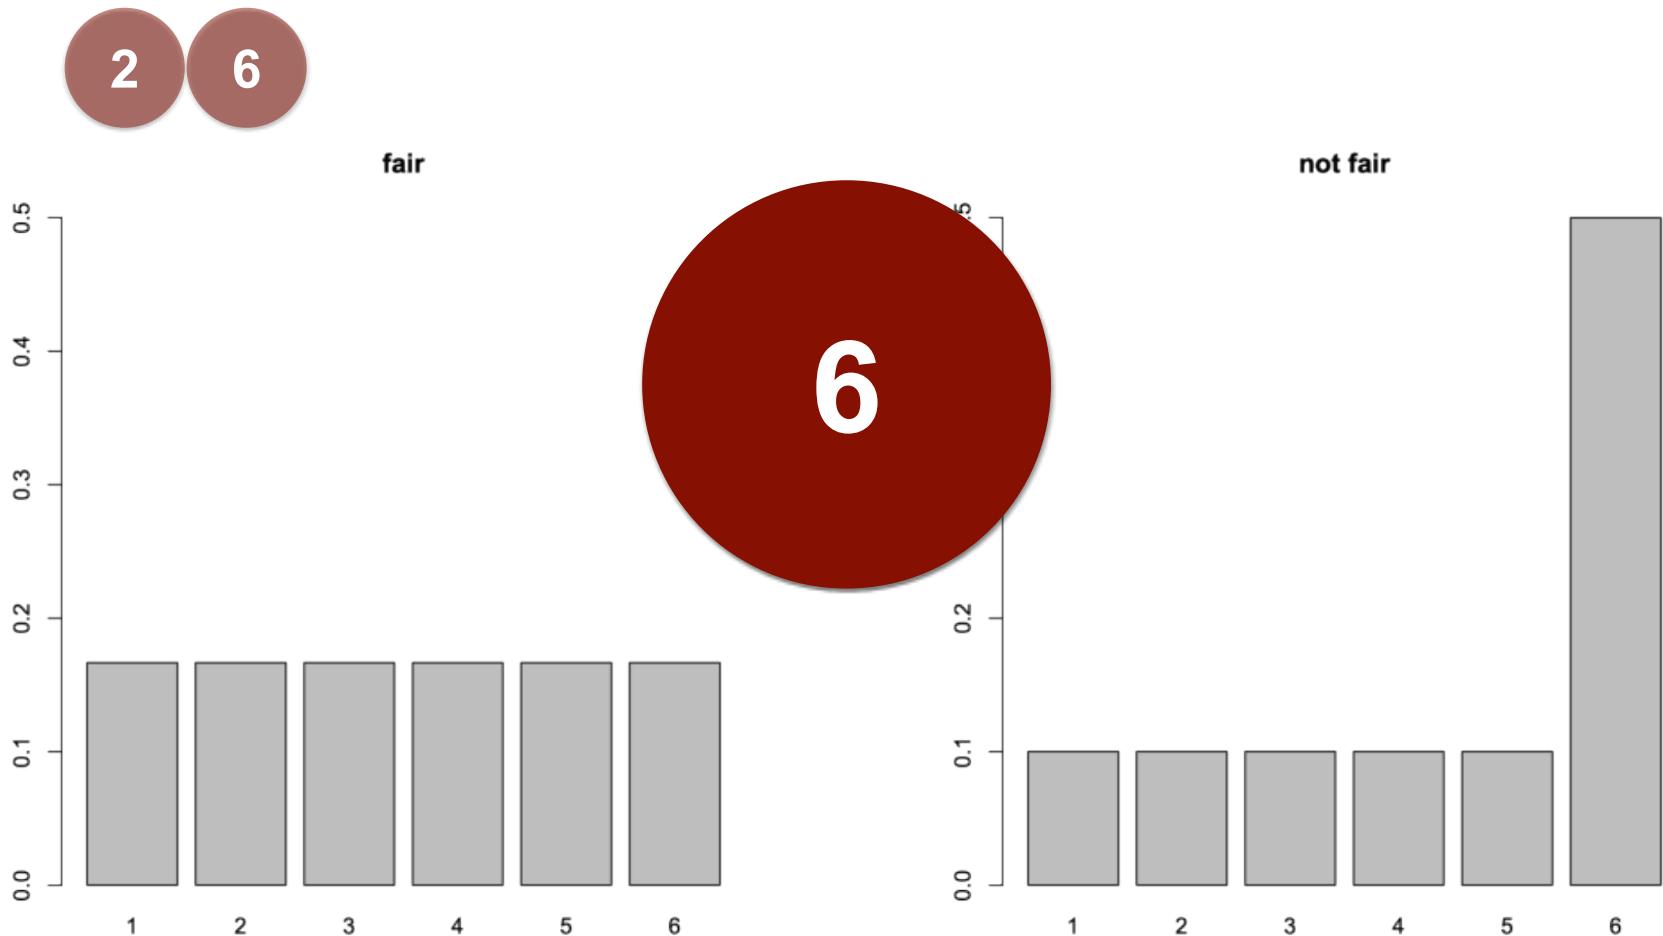
Probability



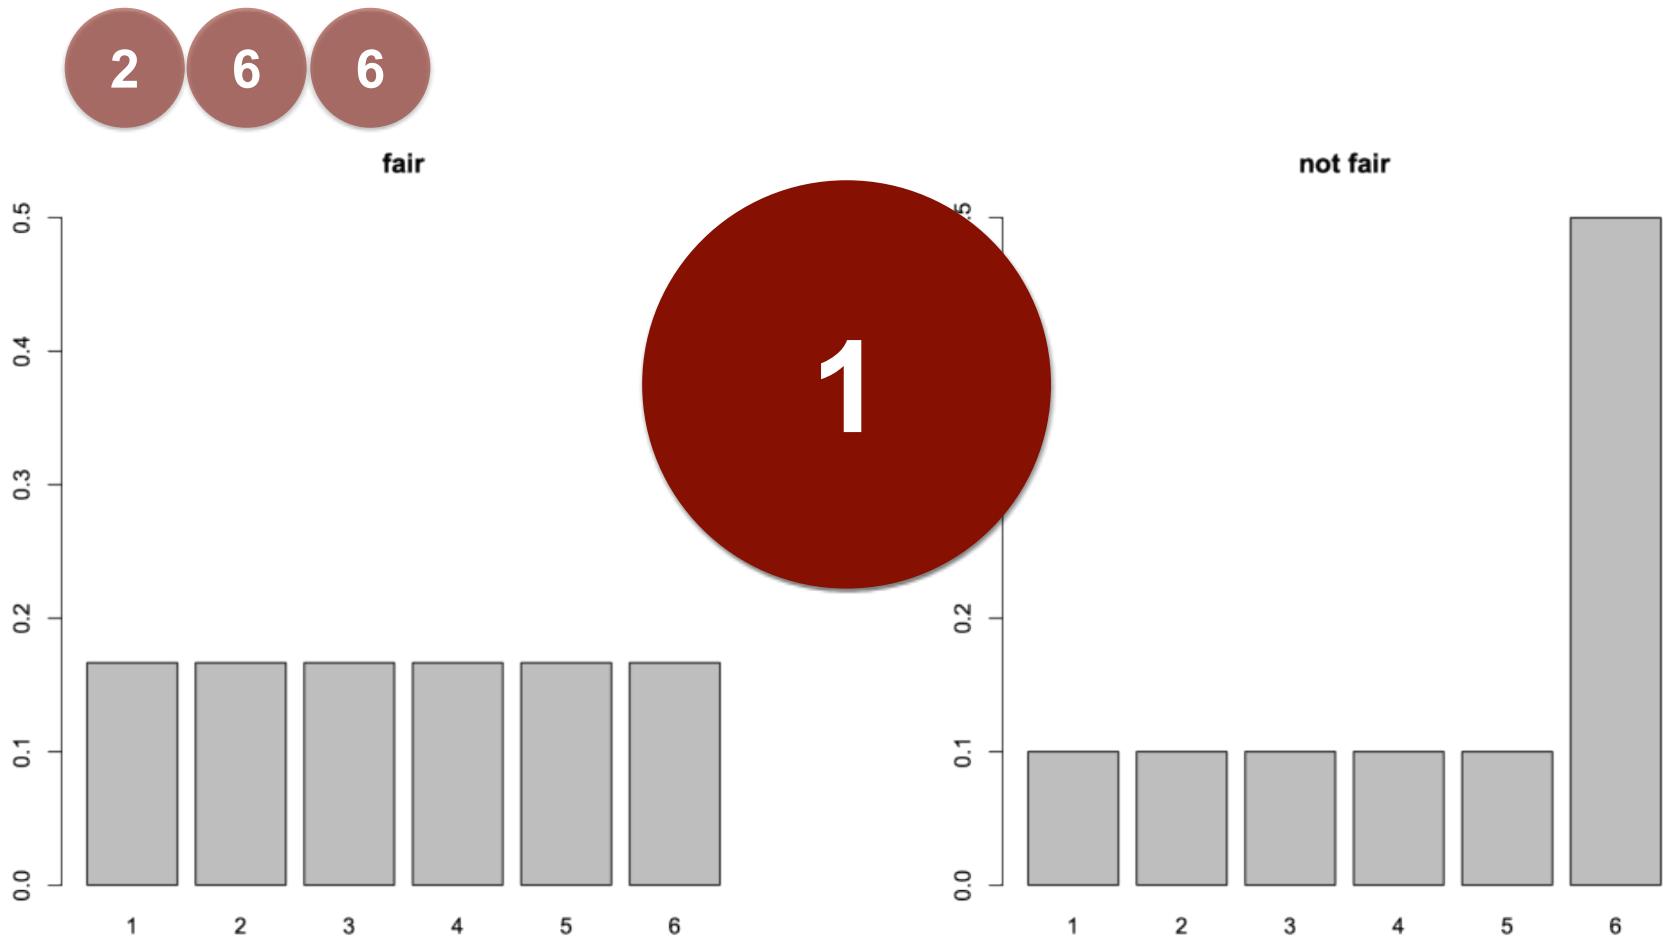
Probability



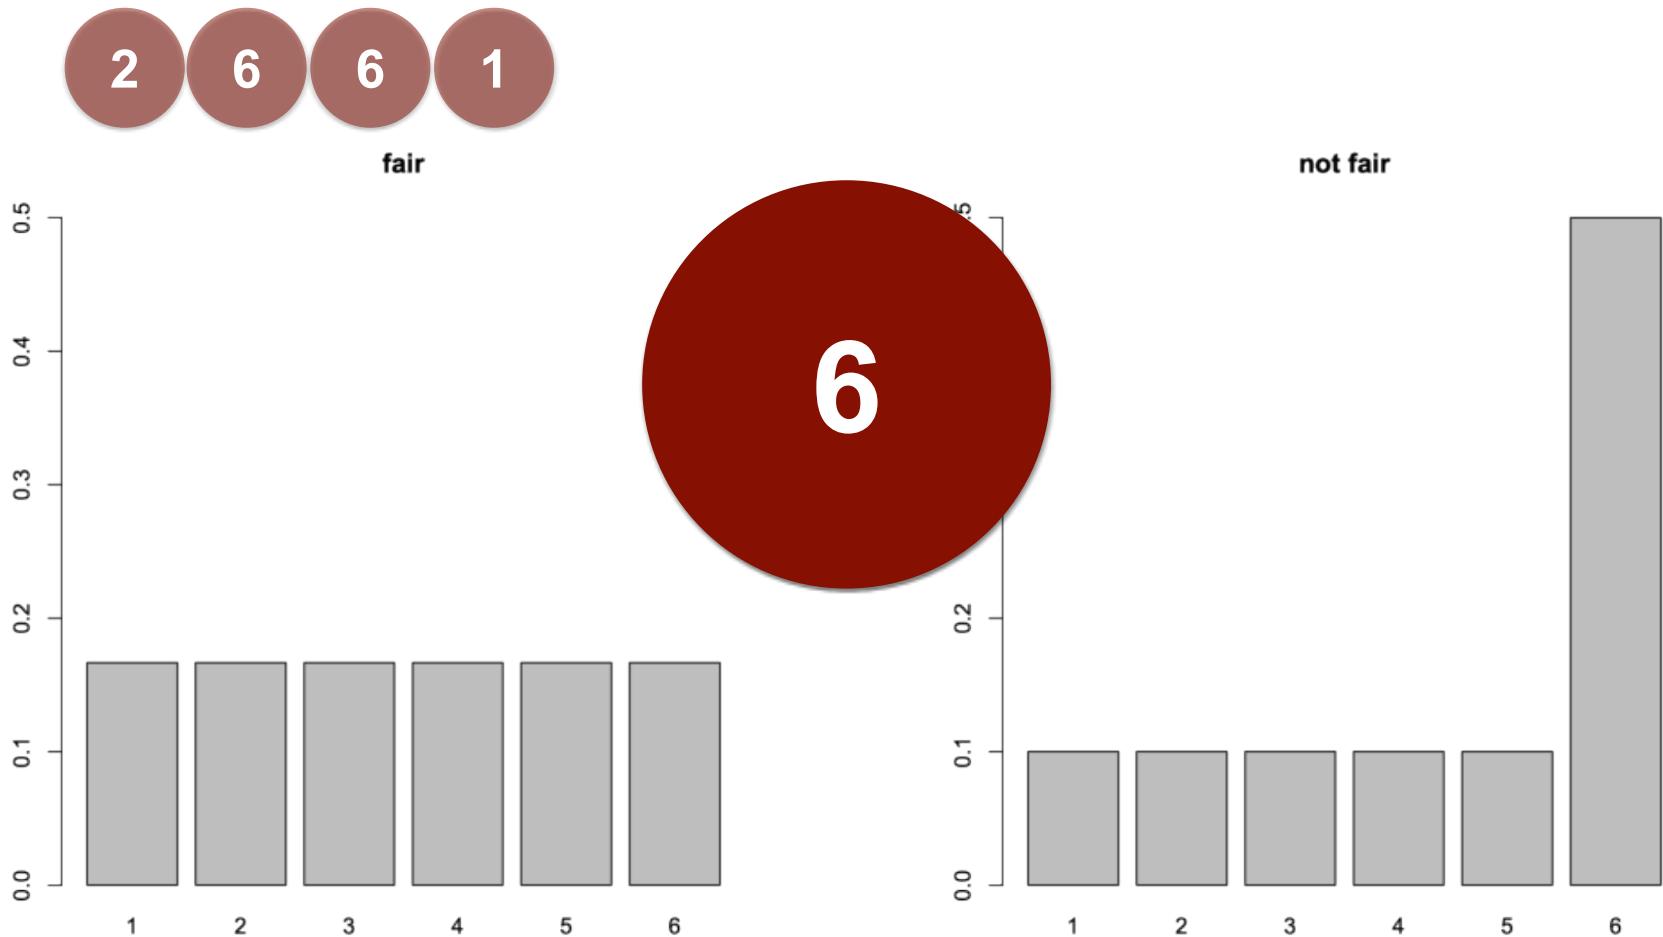
Probability



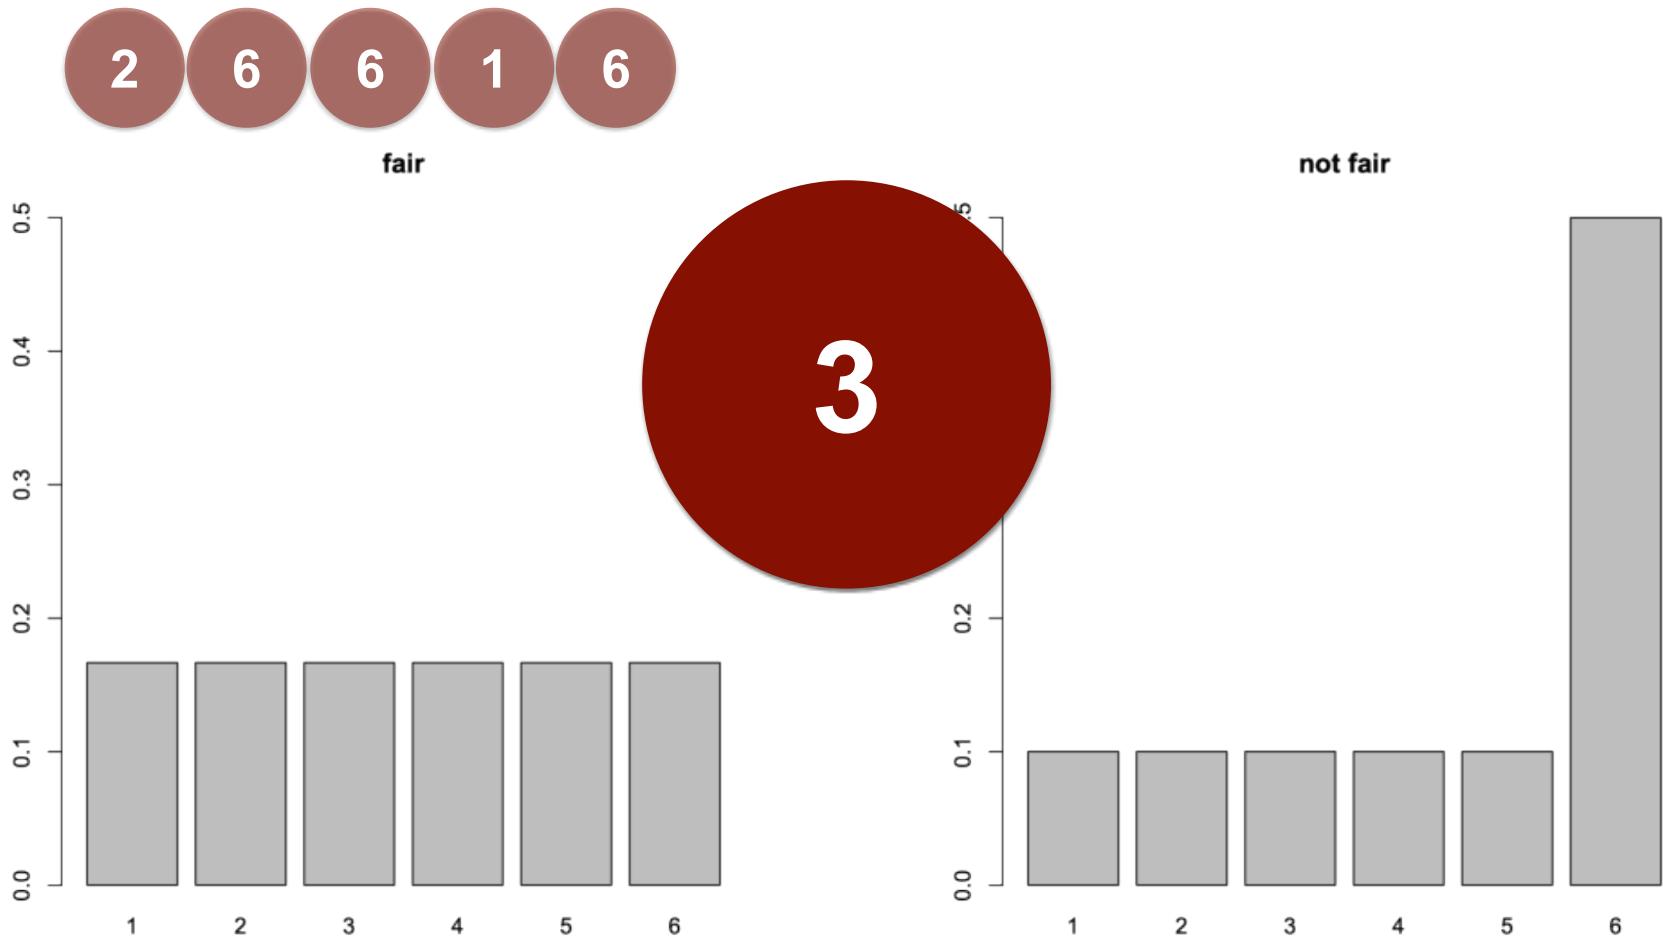
Probability



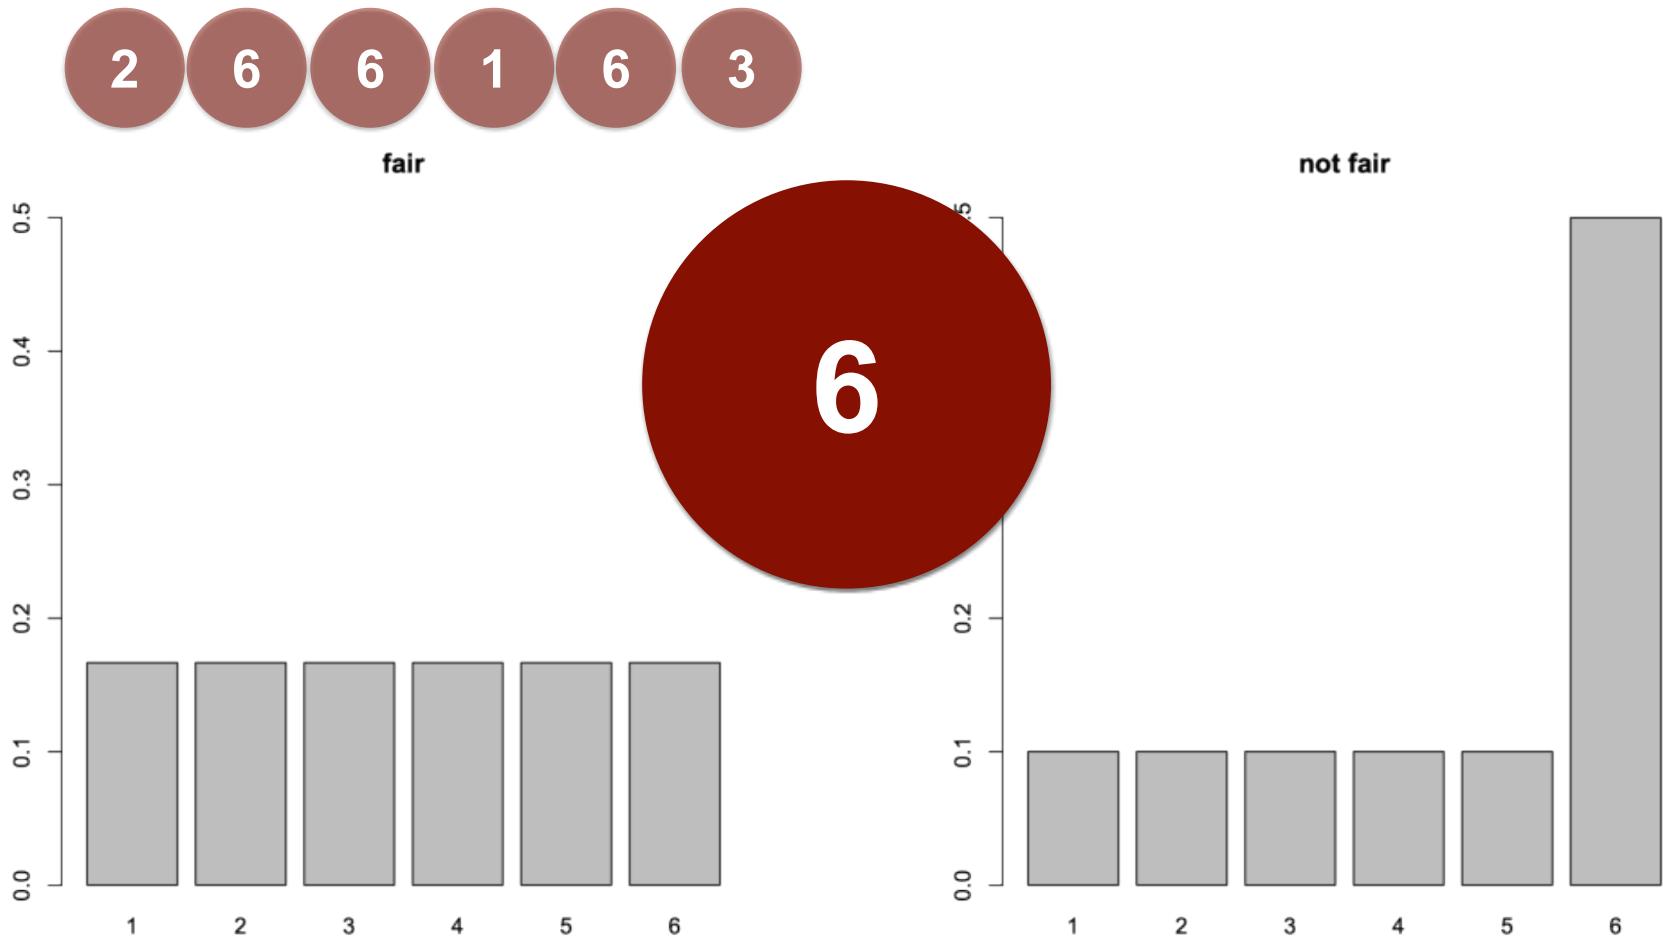
Probability



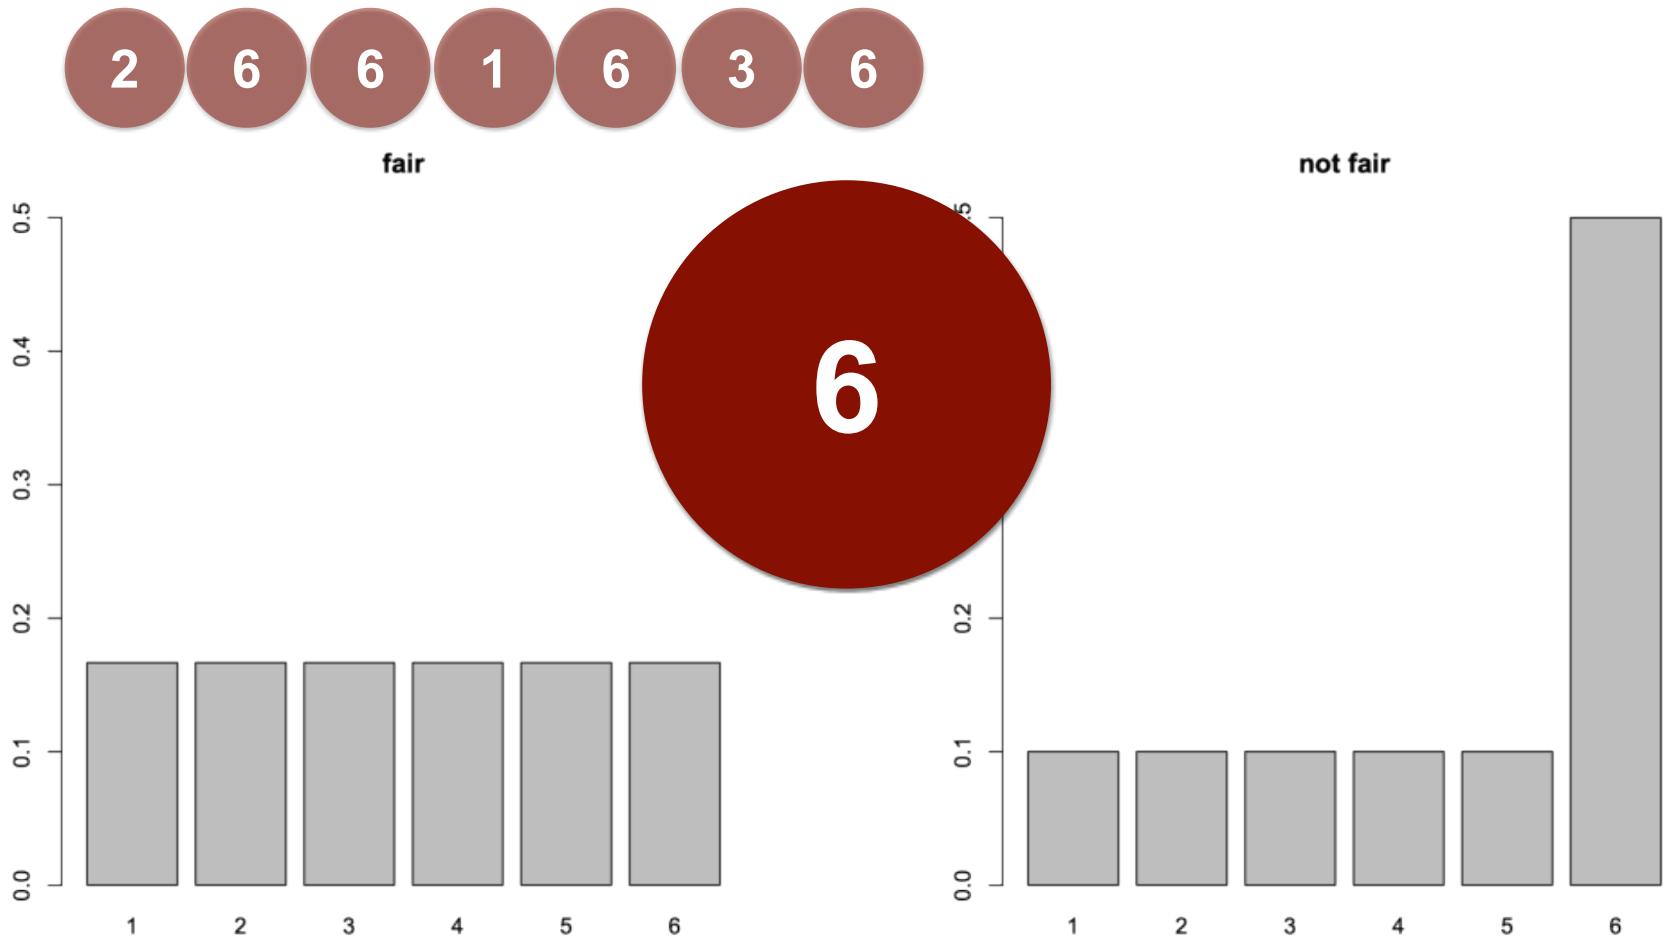
Probability



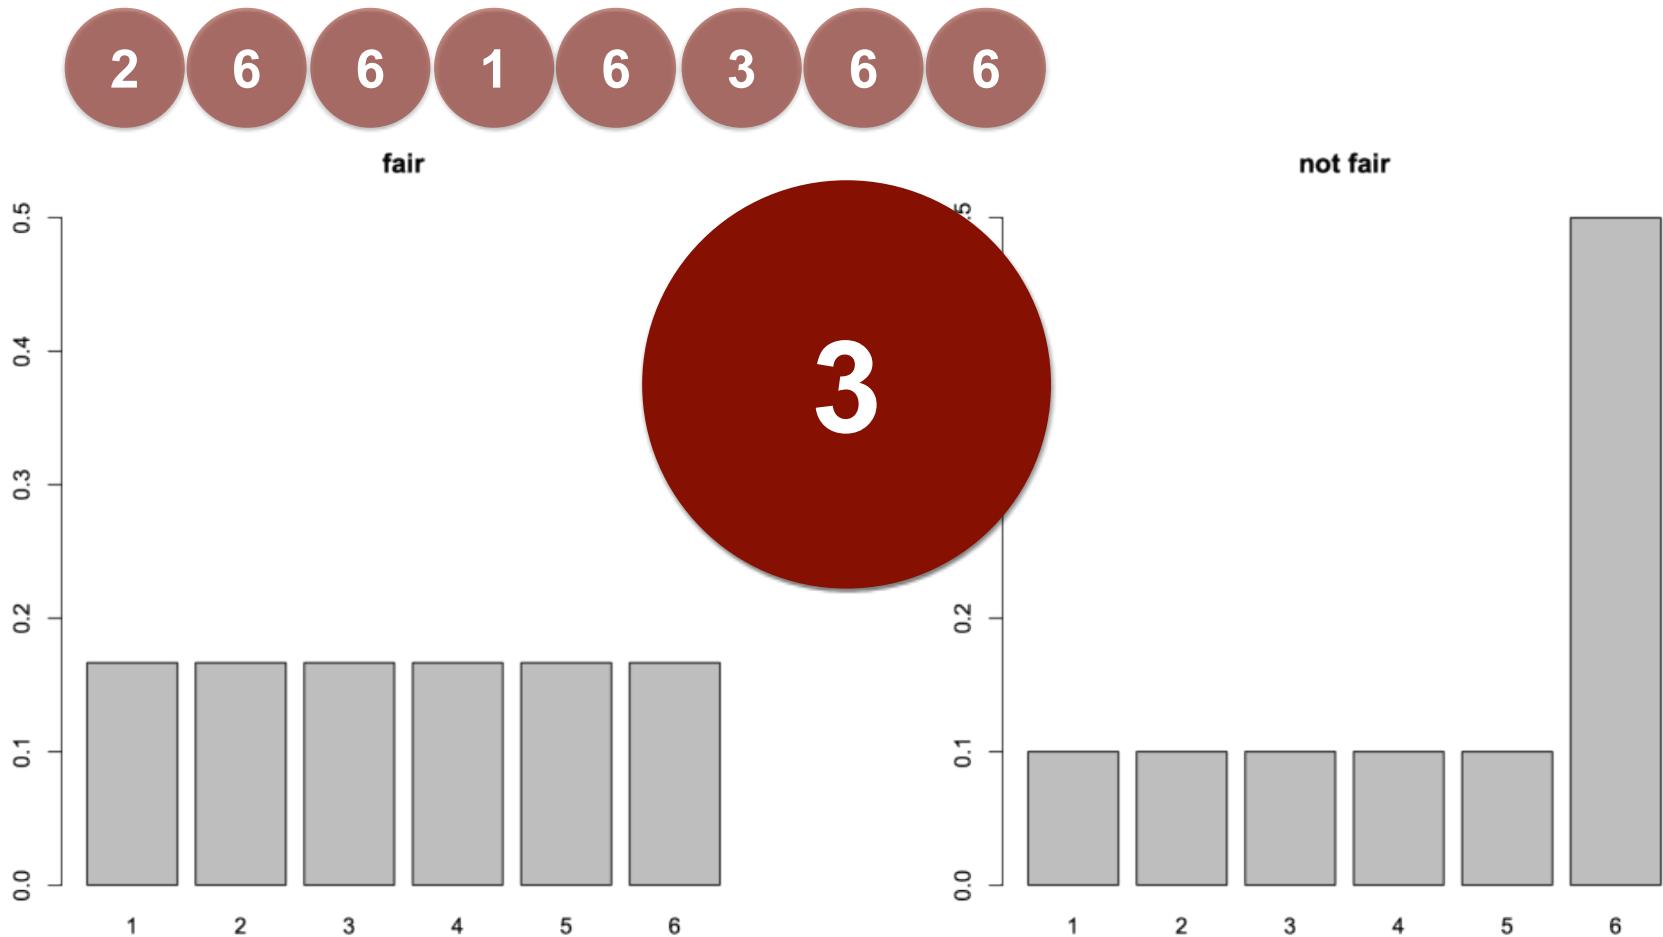
Probability



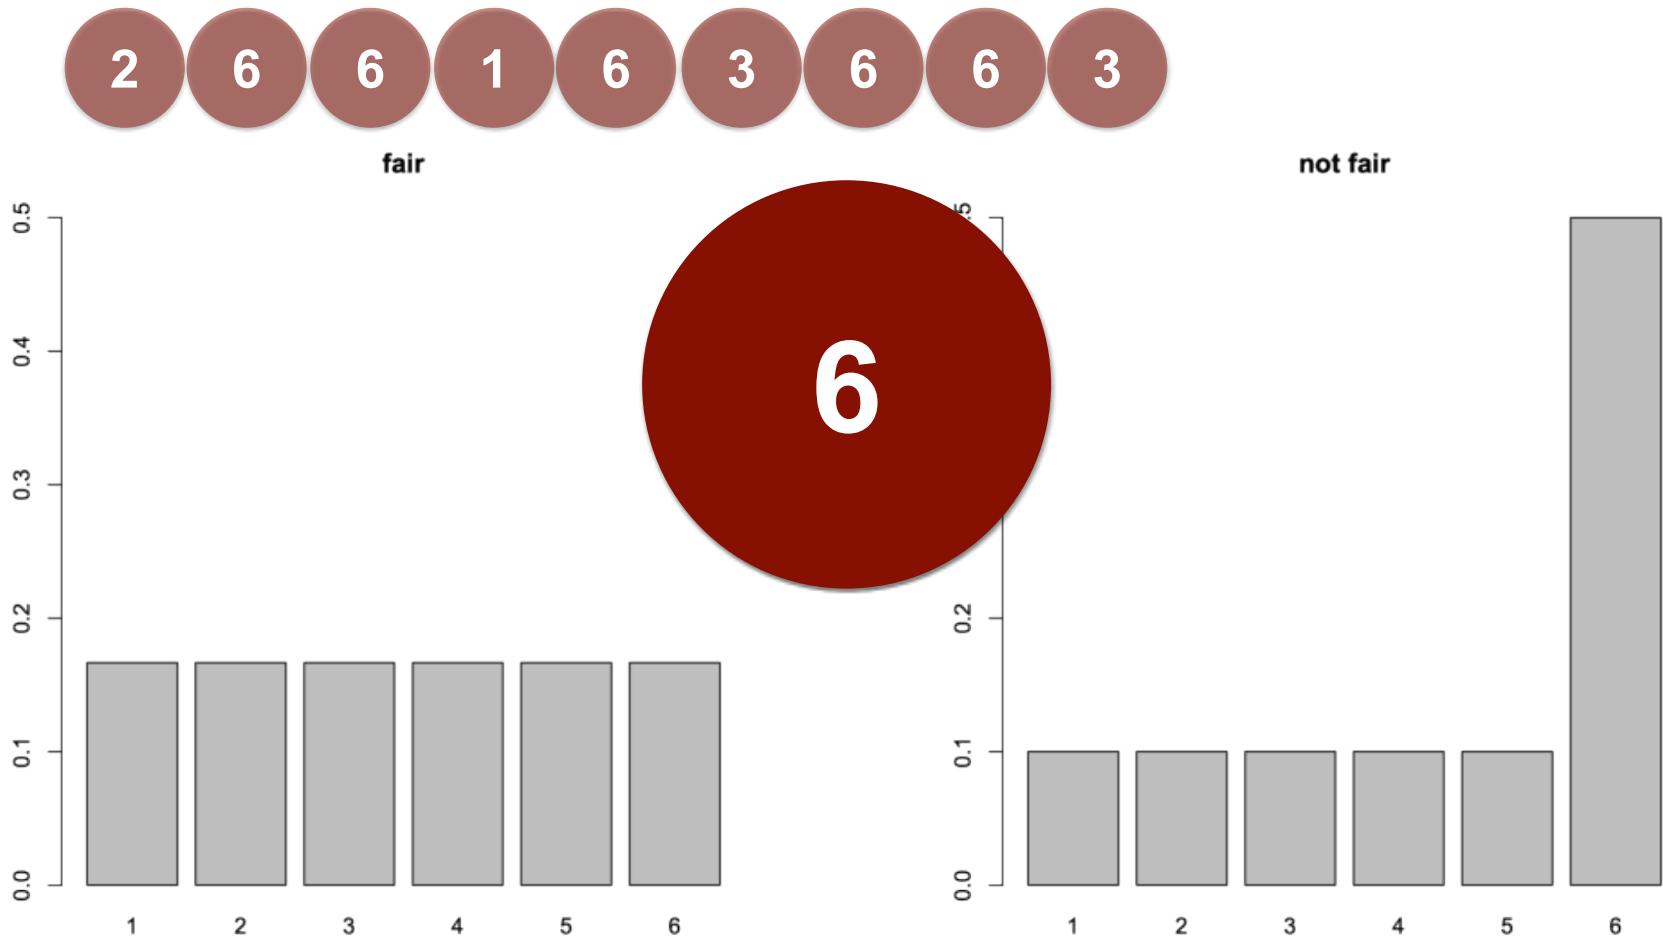
Probability



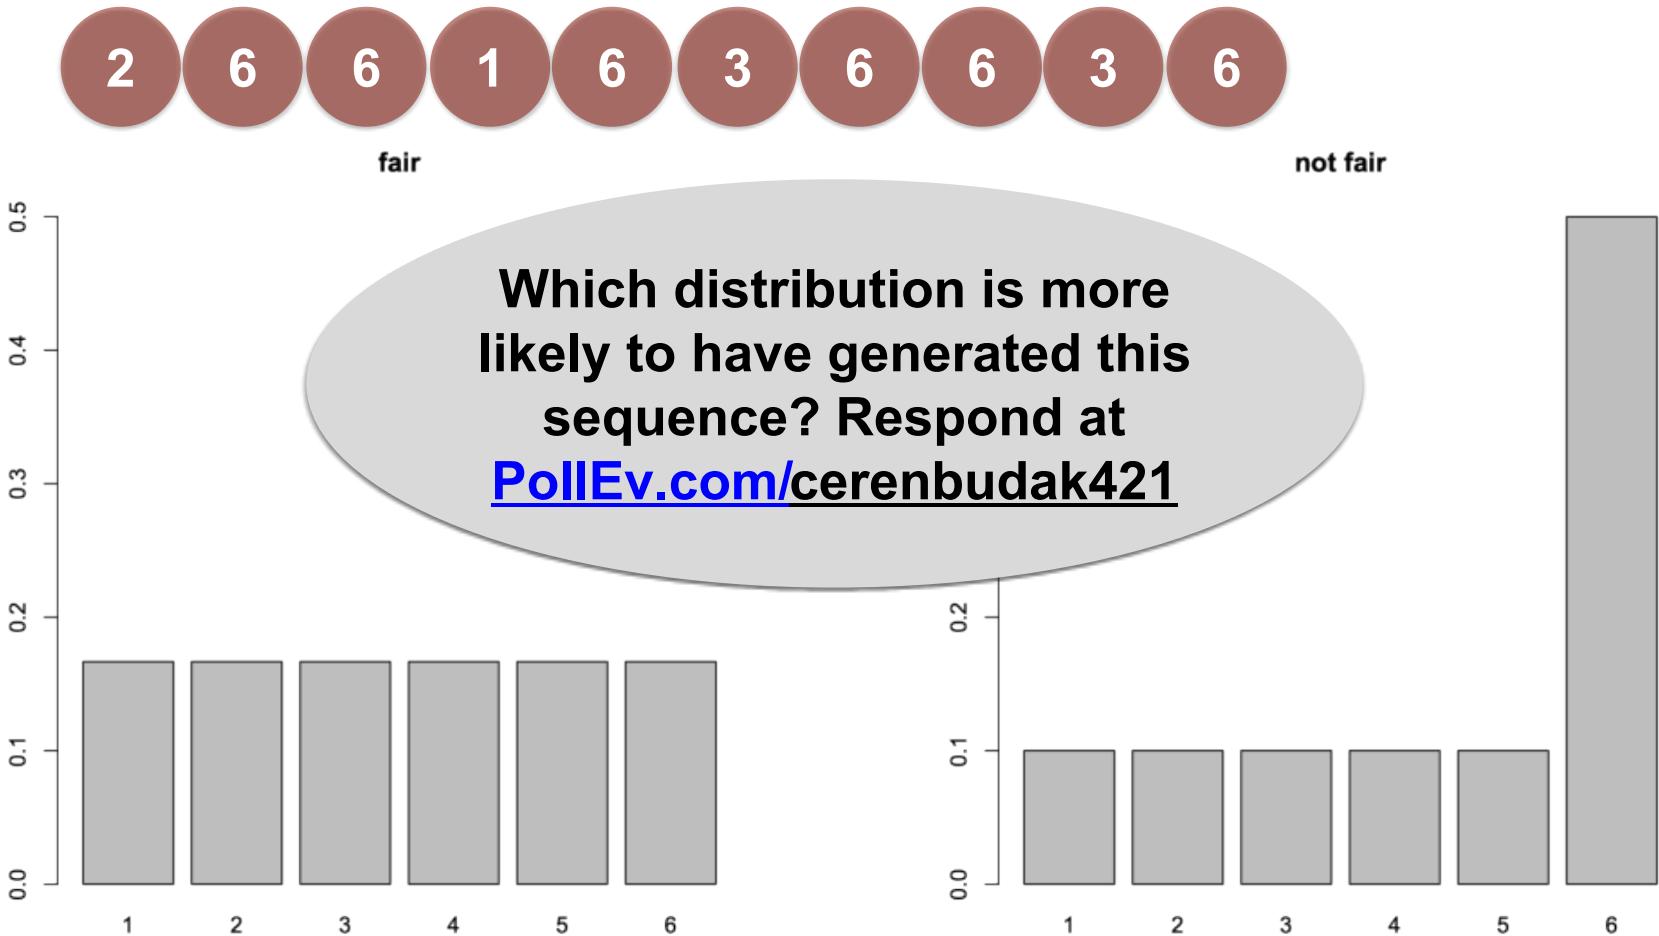
Probability



Probability



Probability





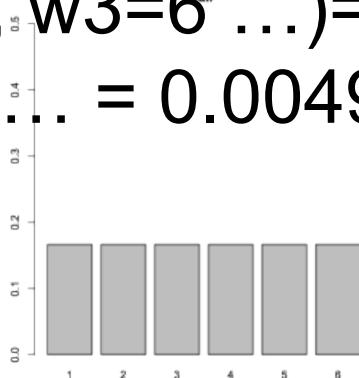
Log in to Poll Everywhere

To present live activities, please log in to your Poll Everywhere account in a separate window.

[Launch log-in window](#)

Data Likelihood

- $P(w_1=2, w_2=6, w_3=6 \dots) = .17 \times .17 \times .17 \dots = 0.004913$

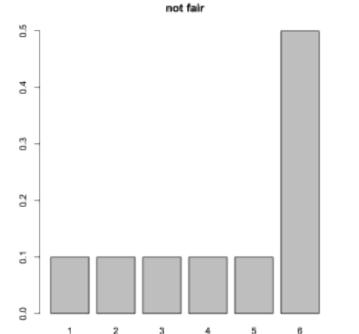
$$P(\text{2} | \text{6} | \text{6})$$


A bar chart with six bars labeled 1 through 6 on the x-axis. The y-axis ranges from 0.0 to 0.5. Each bar has a height of approximately 0.17.

Outcome	Probability
1	~0.17
2	~0.17
3	~0.17
4	~0.17
5	~0.17
6	~0.17

$$=.17 \times .17 \times .17 \\ = 0.004913$$

- $P(w_1=2, w_2=6, w_3=6)$

$$P(\text{2} | \text{6} | \text{6})$$


A bar chart with six bars labeled 1 through 6 on the x-axis. The y-axis ranges from 0.0 to 0.5. The bars have heights of approximately 0.1, 0.1, 0.1, 0.1, 0.1, and 0.5 respectively. The bar for outcome 6 is significantly taller than the others.

Outcome	Probability
1	~0.1
2	~0.1
3	~0.1
4	~0.1
5	~0.1
6	~0.5

$$=.1 \times .5 \times .5 \\ = 0.025$$

Data Likelihood

- The likelihood gives us a way of discriminating between possible alternative parameters, but also a strategy for picking a single best* parameter (θ) among all possibilities

Word choice as weighted dice

Unigram probability

A document
with positive
reviews (or can
be a collection
for
classification)

A document
with negative
reviews
(or can be a
collection for
classification)

Maximum Likelihood Estimate

- This is a maximum likelihood estimate for $P(X)$; the parameter values for which the data we observe (X) is **most likely**.

$$P(X = \text{the}) = \frac{\#\text{the}}{\#\text{total words}}$$

Conditional Probability

- Probability that **one random variable** takes a particular value *given* the fact that **a different variable** takes another

$$P(X = x | Y = y)$$

$$P(X_i = \text{hate} | Y = \oplus)$$

Maximum Likelihood vs. Bayesian

- Maximum likelihood estimation
 - “Best” means “data likelihood reaches maximum”
$$\hat{\theta} = \operatorname{argmax}_{\theta} P(X|\theta)$$
 - Problem: small sample
- Bayesian estimation
 - “Best” means being consistent with our “prior” knowledge and explaining data well
 - Problem: how to define prior?

2. Maximum Likelihood (MLE)

- **目标:** 选择能让观测数据概率最大的参数值 $\hat{\theta}$ 。
- **做法:**

$$\hat{\theta}_{MLE} = \arg \max_{\theta} P(\text{data}|\theta)$$

- **特点:**

- 只依赖 观测到的数据。
- 不考虑任何先验知识 (prior)。
- 通常给出一个点估计 (single best estimate)。
- 简单、直接，但容易过拟合。

3. Bayesian Estimation

- **目标:** 求出参数的后验分布 $P(\theta|\text{data})$ ，然后用它做决策。
- **公式:**

$$P(\theta|\text{data}) = \frac{P(\text{data}|\theta)P(\theta)}{P(\text{data})}$$

其中：

- $P(\theta)$: 先验概率 (prior)
- $P(\text{data}|\theta)$: 似然 (likelihood)
- $P(\theta|\text{data})$: 后验概率 (posterior)

- **特点:**

- 同时利用 观测数据 和 先验知识。
- 得到的不是单一数值，而是整个概率分布。
- 更鲁棒，尤其是在数据量少的情况下。

Probabilistic Approaches to Text Retrieval

Probabilistic Approach to IR

- Document ranking based on the **conditional probability of relevance**

$$P(R = 1 | D, Q)$$

Relevance Document Query
 $R \in \{0, 1\}$

Ranking based on

$$\text{Odds-ratio} \rightarrow O(R | D, Q) = \frac{P(R=1 | D, Q)}{P(R=0 | D, Q)}$$

Problem: Hard to compute $P(R|D, Q)$ directly

Odds (几率)

几率是“发生的概率”和“不发生的概率”的比值：

$$\text{odds} = \frac{P(\text{事件发生})}{P(\text{事件不发生})}$$

比如：

- 事件发生概率 $P = 0.8 \rightarrow \text{几率} = 0.8/0.2 = 4$
- 事件发生概率 $P = 0.5 \rightarrow \text{几率} = 0.5/0.5 = 1$
- 事件发生概率 $P = 0.2 \rightarrow \text{几率} = 0.2/0.8 = 0.25$

几率越大，表示事件更有可能发生。

Odds Ratio (几率比)

几率比是两个几率的比值。

在 IR 中，我们用的是相关和不相关的几率比：

$$\text{odds-ratio} = \frac{\text{相关的几率}}{\text{不相关的几率}} = \frac{P(R = 1|D, Q)}{P(R = 0|D, Q)}$$

当这个值 $> 1 \rightarrow$ 文档更可能是相关的；

当这个值 $< 1 \rightarrow$ 文档更可能是不相关的。

为什么用odd:

1. 数学上更方便
2. 区分度更大

Probabilistic Retrieval Models: Computing $p(R|Q,D)$

- Basic idea
 - Compute $P(R|Q,D)$ using Bayes' rule

$$P(R|Q,D) = \frac{P(Q,D|R)P(R)}{P(Q,D)}$$

- $P(D,Q|R=r)$ is the probability that if we know that a relevant document is retrieved, it's D (for query Q)
 - $P(R=r)$ is the probability of retrieving a relevant document
 - $P(D,Q)$ probability of retrieving D and issuing Q
 - Assumption: computing $P(Q, D|R)$ is easier.
 - But what about $P(R)$ and $P(Q,D)$?
-
- 对所有文档, $p(Q|D)$ 是常数 → 排序时可以忽略
 - 如果假设相关性 R 和文档内容 D 无关 (先验相同) , $p(R|D)$ 也是常数 → 排序时也可以忽略

Probabilistic Retrieval Models: Computing $p(R|Q,D)$

- Assumption: computing $P(Q, D|R)$ is easier.

$$\text{Odds ratio } \alpha(R|Q,D) = \frac{P(R=1|Q,D)}{P(R=0|Q,D)} = \frac{P(Q,D|R=1)P(R=1)/P(Q,D)}{P(Q,D|R=0)P(R=0)/P(Q,D)} = \frac{P(Q,D|R=1)}{P(Q,D|R=0)} \frac{\frac{P(R=1)}{P(R=0)}}{\frac{P(R=1)}{P(R=0)}}$$

↑
Ignored for ranking D.
Why?

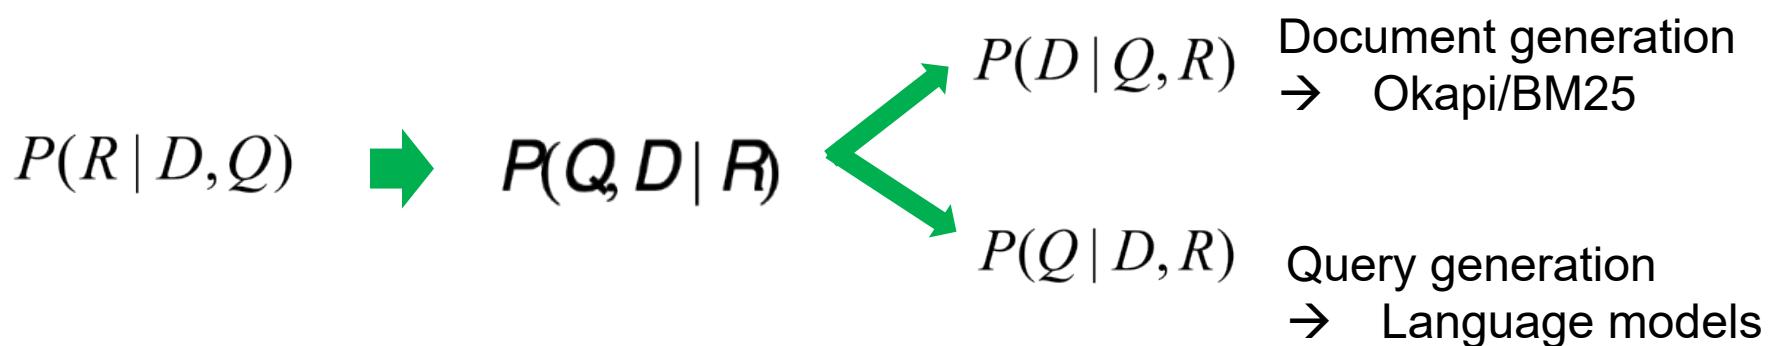
- How to compute $P(Q,D|R)$?
 - Document “generation”: $P(Q,D|R) = P(D|Q,R) P(Q|R)$
考虑给定查询的情况下，文档 D 被“生成”的概率
 - Query “generation”: $P(Q,D|R) = P(Q|D,R) P(D|R)$
考虑给定文档的情况下，查询 Q 被“生成”的概率

Query Generation and Document Generation

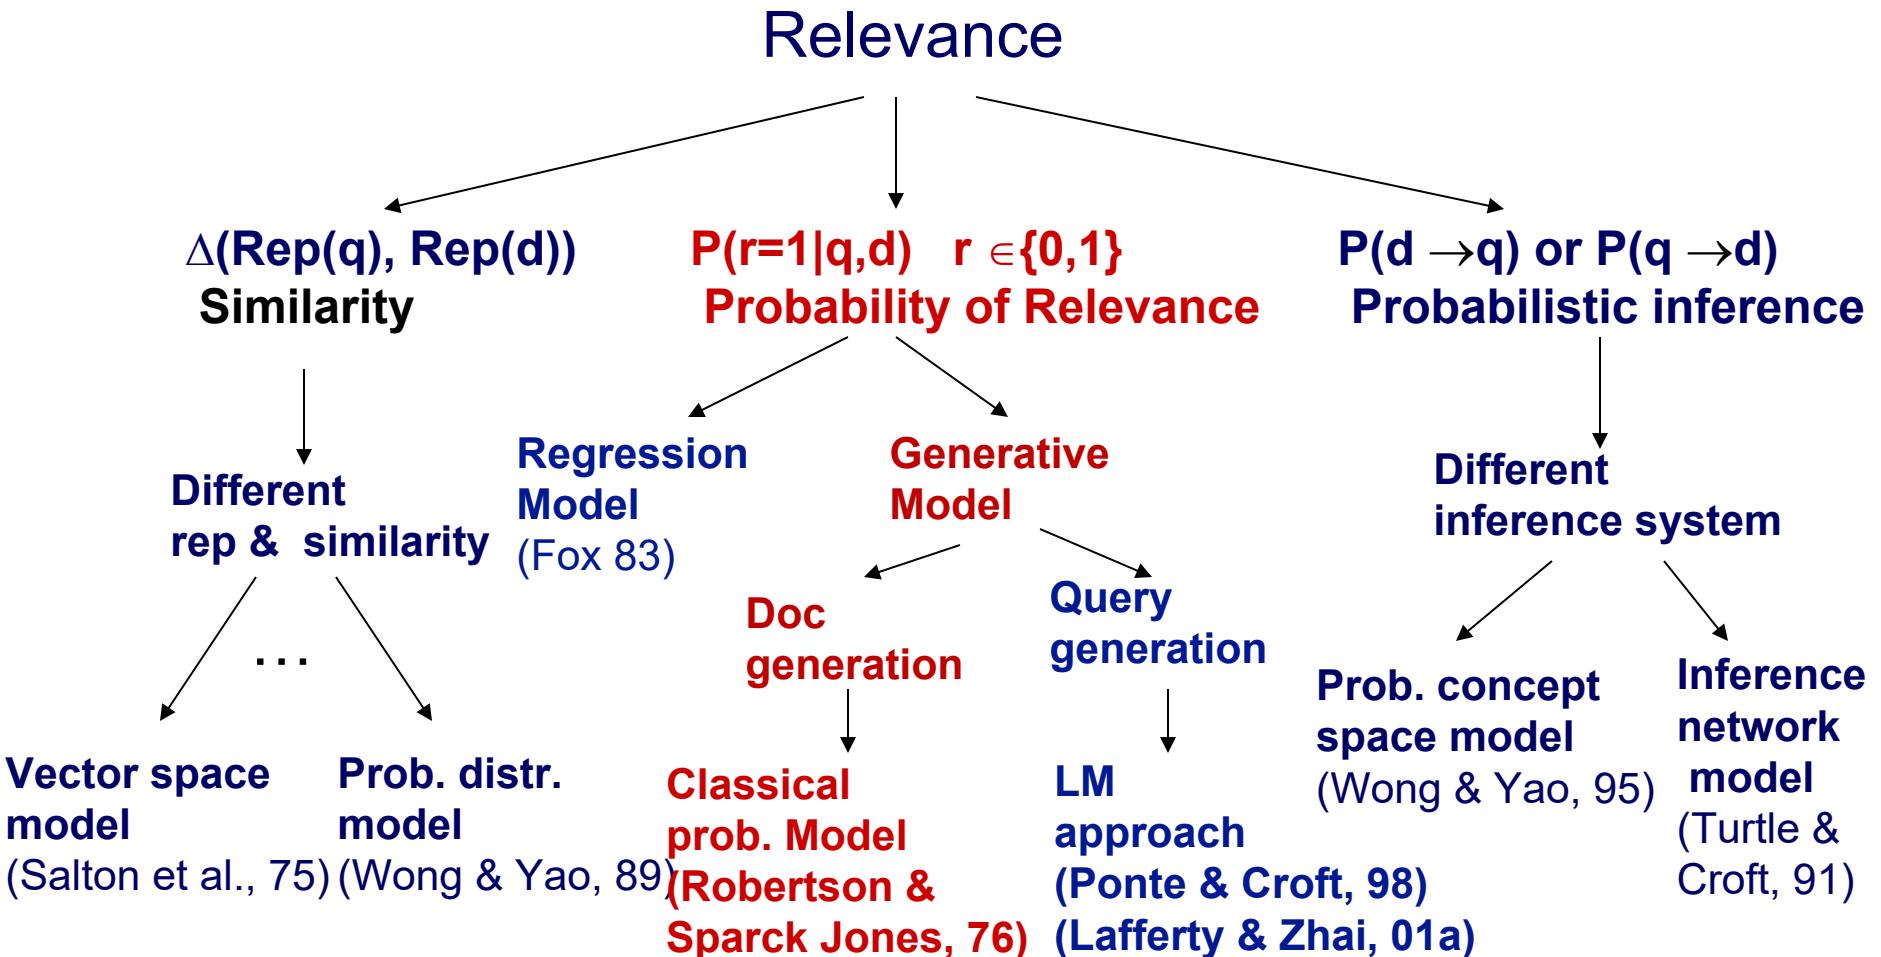
Ranking based on

假设文档 D 确实是相关的 ($R=1$)
然后计算“用户发出了 Q 这个查询，且我们看到了 D
这篇文档”同时发生的概率

$$O(D, Q) = \frac{P(R=1 | D, Q)}{P(R=0 | D, Q)} \sim \frac{P(Q, D | R=1)}{P(Q, D | R=0)}$$



The Notion of Relevance



Binary Independence Retrieval Model

- Assumes that query and doc terms are independent: it's a **Naive Bayes Classifier** for IR
- Documents are represented by **binary vectors**
 - if a term is present in a document, it's 1, otherwise it's 0

$$score(Q, D) = \log \frac{P(R = r | D, Q)}{P(R = \neg r | D, Q)}$$

We're not going to produce a probability here. why?

the log-odds ratio of being relevant or not

Binary Independence Retrieval Model

Bayes Rule

$$score(Q, D) = \log \frac{P(R = r | D, Q)}{P(R = \neg r | D, Q)} = \log \frac{P(D, Q | R = r)P(R = r)}{P(D, Q | R = \neg r)P(R = \neg r)}$$

$P(R=r)$ and $P(R=\neg r)$ are just constants and will not change relative positions of documents in the rating, so let's remove them:

$$score(Q, D) = \log \frac{P(D, Q | R = r)}{P(D, Q | R = \neg r)}$$

Binary Independence Retrieval Model

$$score(Q, D) = \log \frac{P(D, Q | R = r)}{P(D, Q | R = \neg r)}$$

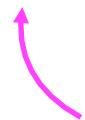
- We assume each term in D is independent, so

$$P(D, Q | R) = \prod_i P(T_i | Q, R)$$

Distribute the product

Log rule for products

$$s(Q, D) = \log \frac{\prod P(T_i | Q, R = r)}{\prod P(T_i | Q, R = \neg r)} = \log \prod \frac{P(T_i | Q, R = r)}{P(T_i | Q, R = \neg r)} = \sum \log \frac{P(T_i | Q, R = r)}{P(T_i | Q, R = \neg r)}$$



T_i is presence of a term T_i in D

Binary Independence Retrieval Model

$$s(Q, D) = \sum_{i \in Q} \log \frac{P(T_i \mid Q, R = r)}{P(T_i \mid Q, R = \neg r)}$$

- Let's define some terms:

$$p_i = P(T_1 = 1 \mid Q, R = r) \quad q_i = P(T_1 = 1 \mid Q, R = \neg r)$$

	$T_i = 1$	p_i	q_i
	$T_i = 0$	$(1-p_i)$	$(1 - q_i)$

- If we assume for all terms not occurring in the query that $p_i = q_i$, we can redefine the score as (a few steps skipped here. Included in additional slides)

$$s(Q, D) = \sum_{i \in Q \cap D} \log \frac{p_i(1 - q_i)}{(1 - p_i)q_i}$$

- $p_i = P(T_i = 1 \mid Q, R = 1)$: 词 w_i 出现在相关文档里的概率
- $q_i = P(T_i = 1 \mid Q, R = 0)$: 词 w_i 出现在不相关文档里的概率

How do we get p_i and q_i ?

$$s(Q, D) = \sum_{i \in Q \cap D} \log \frac{p_i(1 - q_i)}{(1 - p_i)q_i} \quad p_i = P(T_1 = 1 \mid Q, R = r) \\ q_i = P(T_1 = 1 \mid Q, R = \neg r)$$

- given a query Q and a corpus C
 - let N be the number of documents in the collection
 - let R be the number of documents relevant to Q
 - n_i be the # of docs that have term w_i
 - r_i be the # of relevant docs that have term w_i
- N : 语料库中文档总数
 - R : 与查询相关的文档总数
 - n_i : 包含词 w_i 的文档数
 - r_i : 包含词 w_i 且相关的文档数

$$p_i = \frac{r_i + \lambda}{R + 2\lambda} \quad q_i = \frac{n_i - r_i + \lambda}{N - R + 2\lambda}$$

Typically set $\lambda=0.5$ to avoid 0 in log.

Why one alpha for the nominator and two for denominator: intuition is that you add one document that includes and one that doesn't for term t_i

The Okapi/BM25 Retrieval Formula

- Classical probabilistic model (among the best performers!)
- Basic Idea: $P(D, Q|R) = P(D|Q, R)P(Q|R)$ Ignored for ranking
- Many ways to define the distribution $P(D|Q, R)$,
Okapi uses a 2-Poisson model

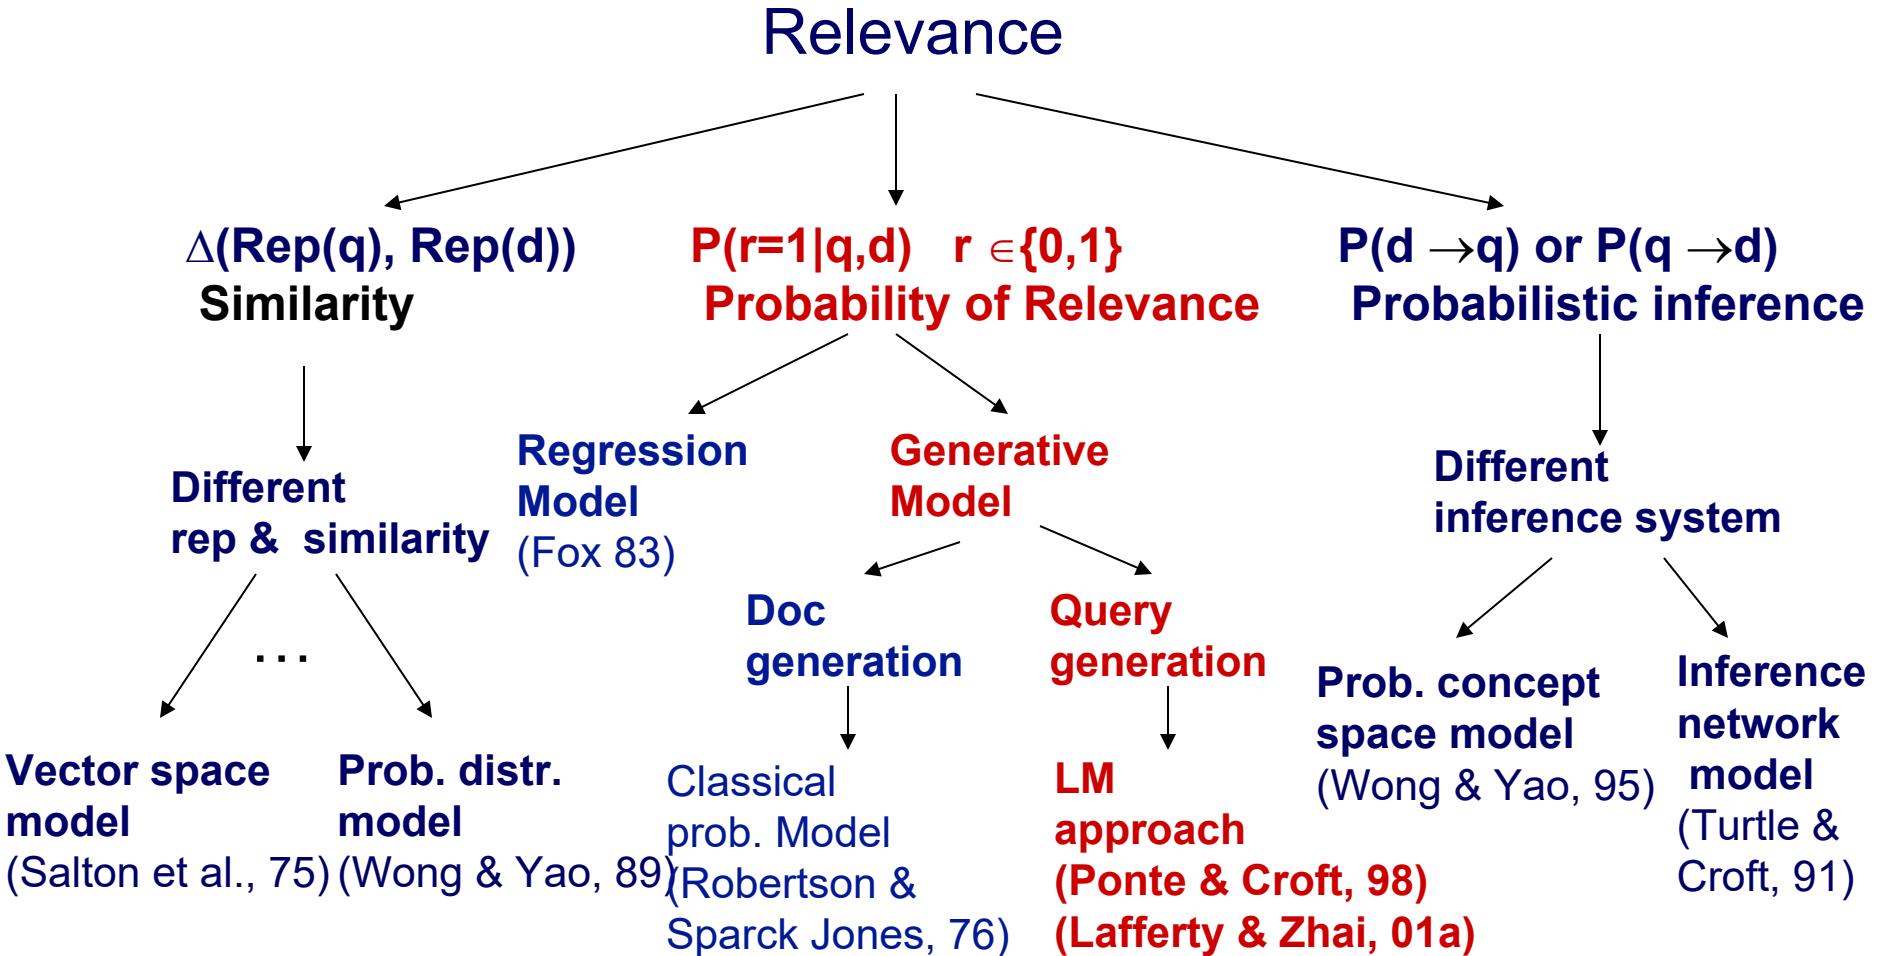
$$Score(d, q) = \sum_{w \in q, w \in d} \left(\log \frac{N - df(w) + 0.5}{df(w) + 0.5} \right) \cdot \frac{(k_1 + 1) \cdot c(w, d)}{k_1((1 - b) + b \frac{|d|}{avdl}) + c(w, d)} \cdot \frac{(k_3 + 1) \cdot c(w, q)}{k_3 + c(w, q)}$$

Only consider words in the query IDF Doc. length Query term frequency Doc. Term frequency

More info for those curious:

https://www.staff.city.ac.uk/~sbrp622/papers/foundations_bm25_review.pdf

The Notion of Relevance



Statistical Language Models

We already know a very simple language model

Word choice as weighted dice

This is a simple language
model

Here, we have *two* language models

positive reviews

negative reviews

Language Model

- Vocabulary \mathcal{V} is a finite set of discrete symbols (e.g., words, characters); $V = |\mathcal{V}|$
- \mathcal{V}^+ is the infinite set of sequences of symbols from \mathcal{V} ; each sequence ends with **STOP**
- $x \in \mathcal{V}^+$

Language Model

$$P(w) = P(w_1, \dots, w_n)$$

$$\begin{aligned} P(\text{"all is well"}) &= \\ P(w_1 = \text{"all"}, w_2 = \text{"is"}, w_3 = \text{"well"}) \times P(\text{STOP}) \end{aligned}$$

$$\sum_{w \in V^+} P(w) = 1 \quad 0 \leq P(w) \leq 1$$

over all sequence lengths!

Language Model

- Language models provide us with a way to quantify the likelihood of sequence of words — i.e., **plausible** sentences.

Language Model

- Language modeling is the task of estimating $P(w)$
- How do we compute this?

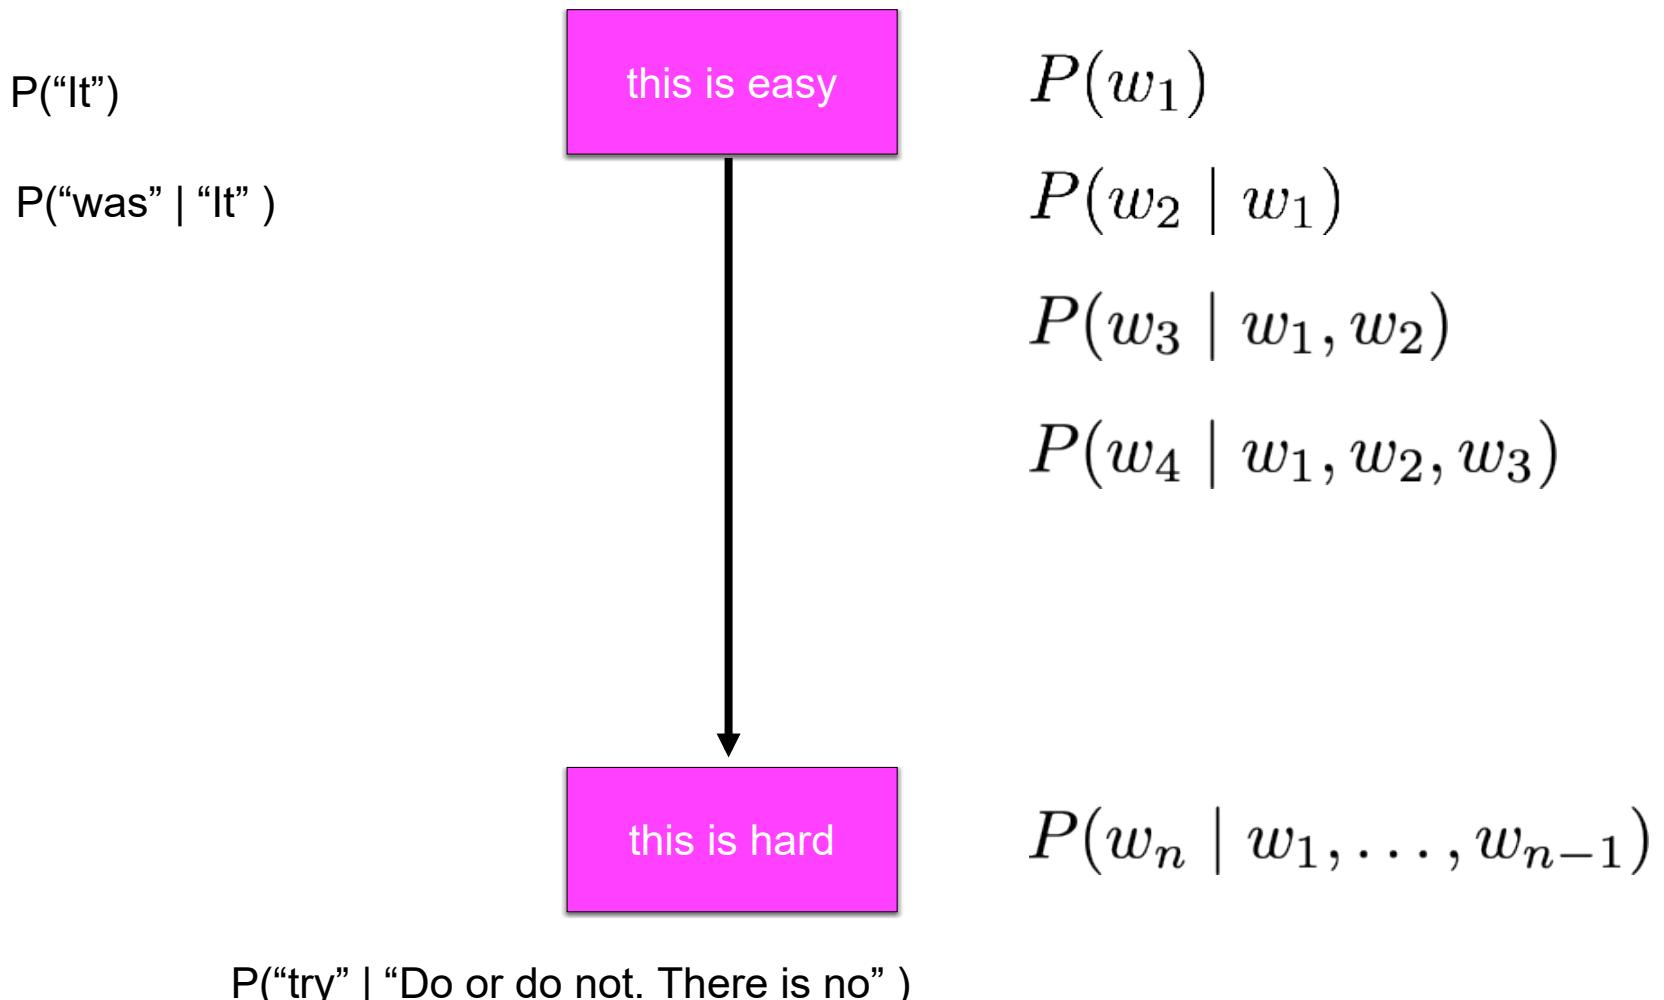
$P(\text{"Do or do not. There is no try."})$



Chain rule (of probability)

$$\begin{aligned} P(x_1, x_2, x_3, x_4, x_5) &= P(x_1) \\ &\quad \times P(x_2 \mid x_1) \\ &\quad \times P(x_3 \mid x_1, x_2) \\ &\quad \times P(x_4 \mid x_1, x_2, x_3) \\ &\quad \times P(x_5 \mid x_1, x_2, x_3, x_4) \end{aligned}$$

Chain rule (of probability)



Markov assumption

first-order

$$P(x_i \mid x_1, \dots x_{i-1}) \approx P(x_i \mid x_{i-1})$$

second-order

$$P(x_i \mid x_1, \dots x_{i-1}) \approx P(x_i \mid x_{i-2}, x_{i-1})$$

Which one of these two give us a bi-gram model?

Respond at PollEv.com/cerenbudak421

First-order



Log in to Poll Everywhere

To present live activities, please log in to your Poll Everywhere account in a separate window.

[Launch log-in window](#)

Estimation

unigram

$$\prod_i^n P(w_i)$$

$$\times P(STOP)$$

bigram

$$\prod_i^n P(w_i \mid w_{i-1})$$

$$\times P(STOP \mid w_n)$$

trigram

$$\prod_i^n P(w_i \mid w_{i-2}, w_{i-1})$$

$$\times P(STOP \mid w_{n-1}, w_n)$$

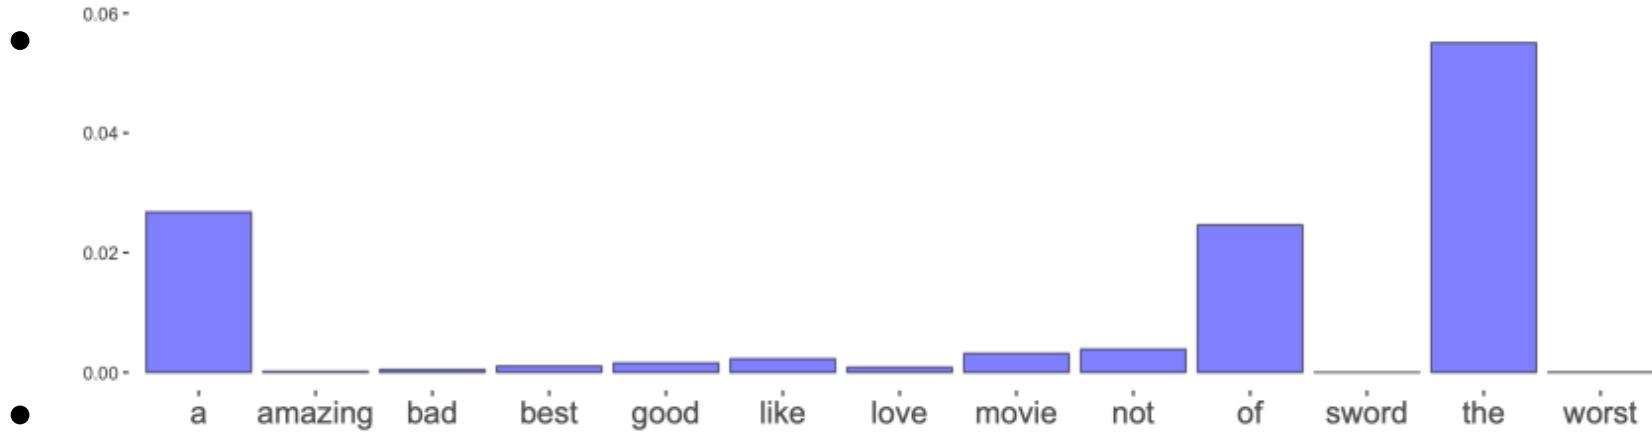
Maximum likelihood estimate

$$\frac{c(w_i)}{N}$$

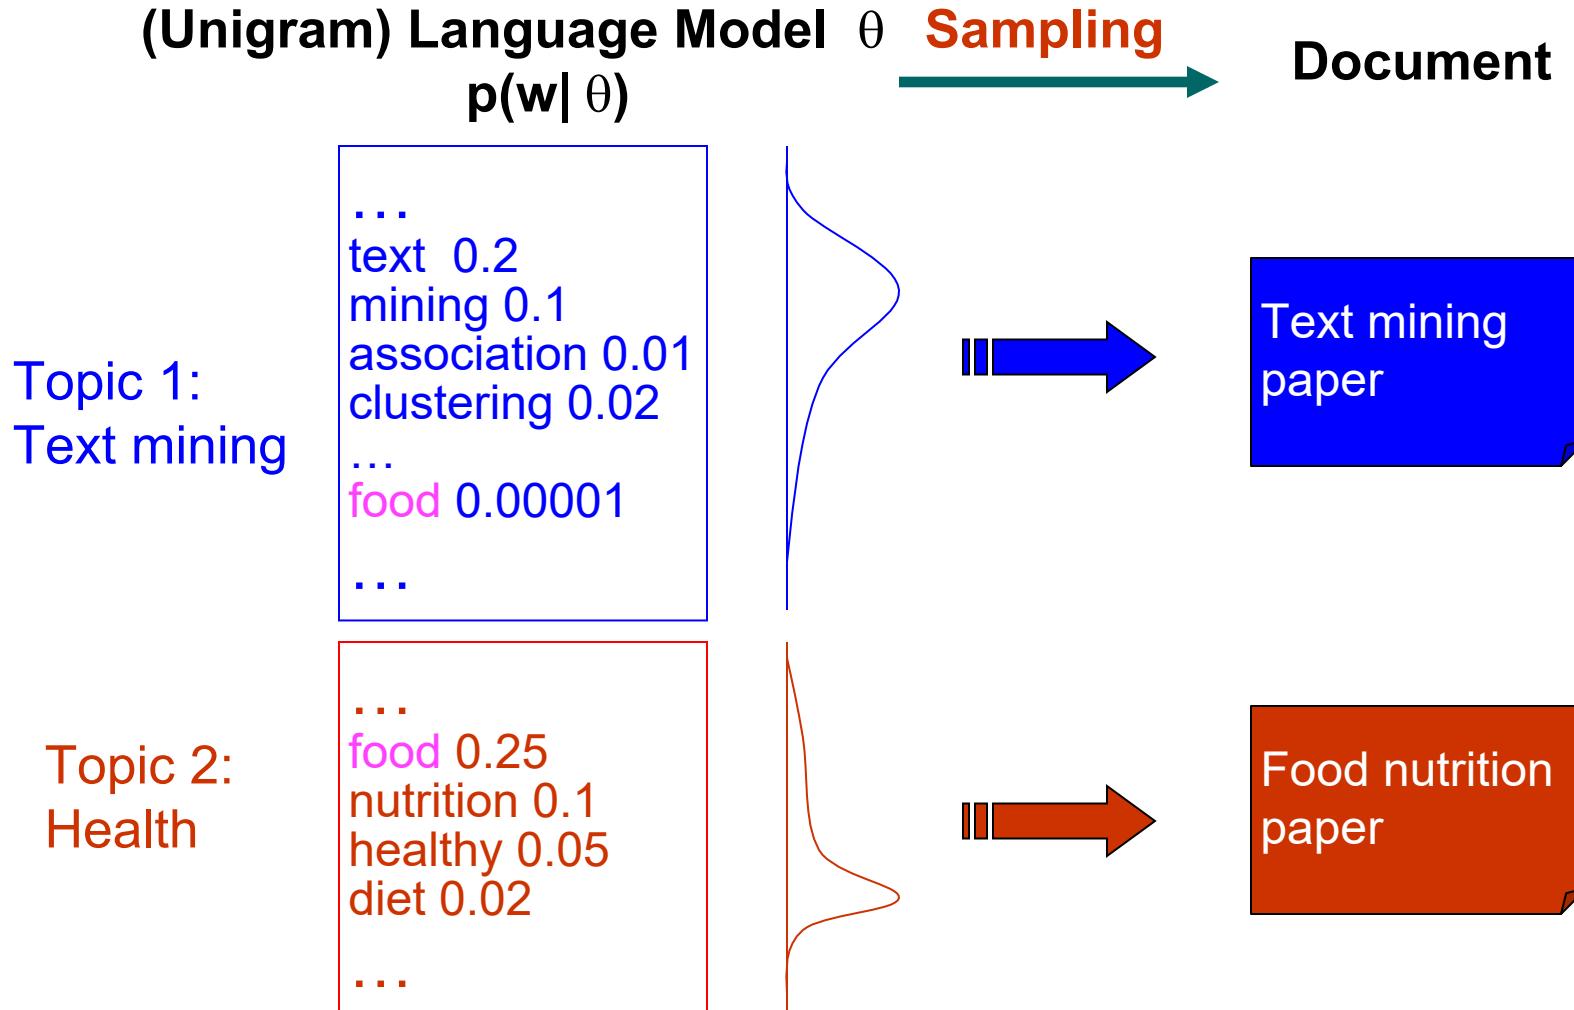
$$\frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

$$\frac{c(w_{i-2}, w_{i-1}, w_i)}{c(w_{i-2}, w_{i-1})}$$

Generating

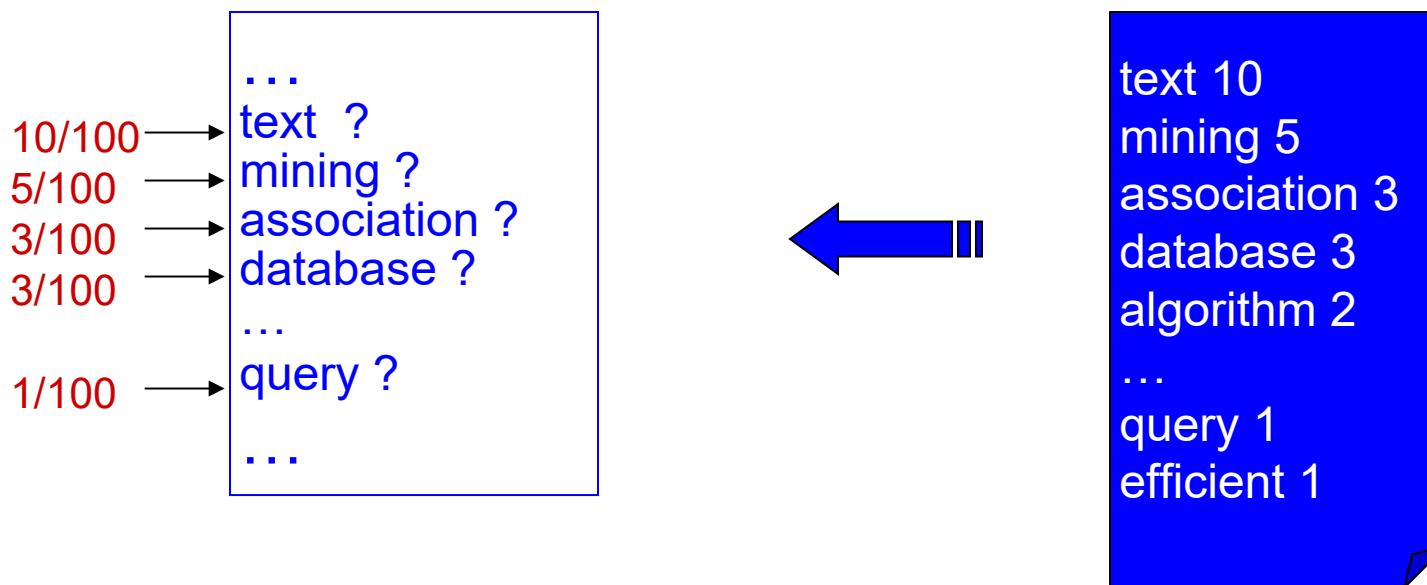


Text Generation with Unigram LM



Estimation of Unigram LM

(Unigram) Language Model θ $p(w|\theta) = ?$ Estimation Document



A “text mining paper”
(total #words=100)

Revisit: Maximum Likelihood Estimate

Data: a document d with counts $c(w_1), \dots, c(w_N)$, and length $|d|$

Model: multinomial distribution M with parameters $\{p(w_i)\}$

Likelihood: $p(d|M)$

Maximum likelihood estimator: $M = \operatorname{argmax}_M p(d|M)$

$$p(w_i) = \frac{c(w_i)}{|d|}$$

Problem: in a short document, if $c(w_i) = 0$, does this mean $p(w_i) = 0$?

Let's go back to our simple unigram model

unigram

$$\prod_i^n P(w_i) \times P(STOP)$$

- What is the probability if one of the words does not occur in the document?
- Solution?

We can combine language models through interpolation (插值)

- A linear interpolation of any two language models p and q (with $\lambda \in [0, 1]$) is also a valid language model.
- Often useful if we have some “background” language model with a much larger vocabulary

$$\lambda p + (1 - \lambda)q$$

p = the web

q = political speeches

Interpolation

$$\lambda p + (1 - \lambda)q$$

p = Star Wars text

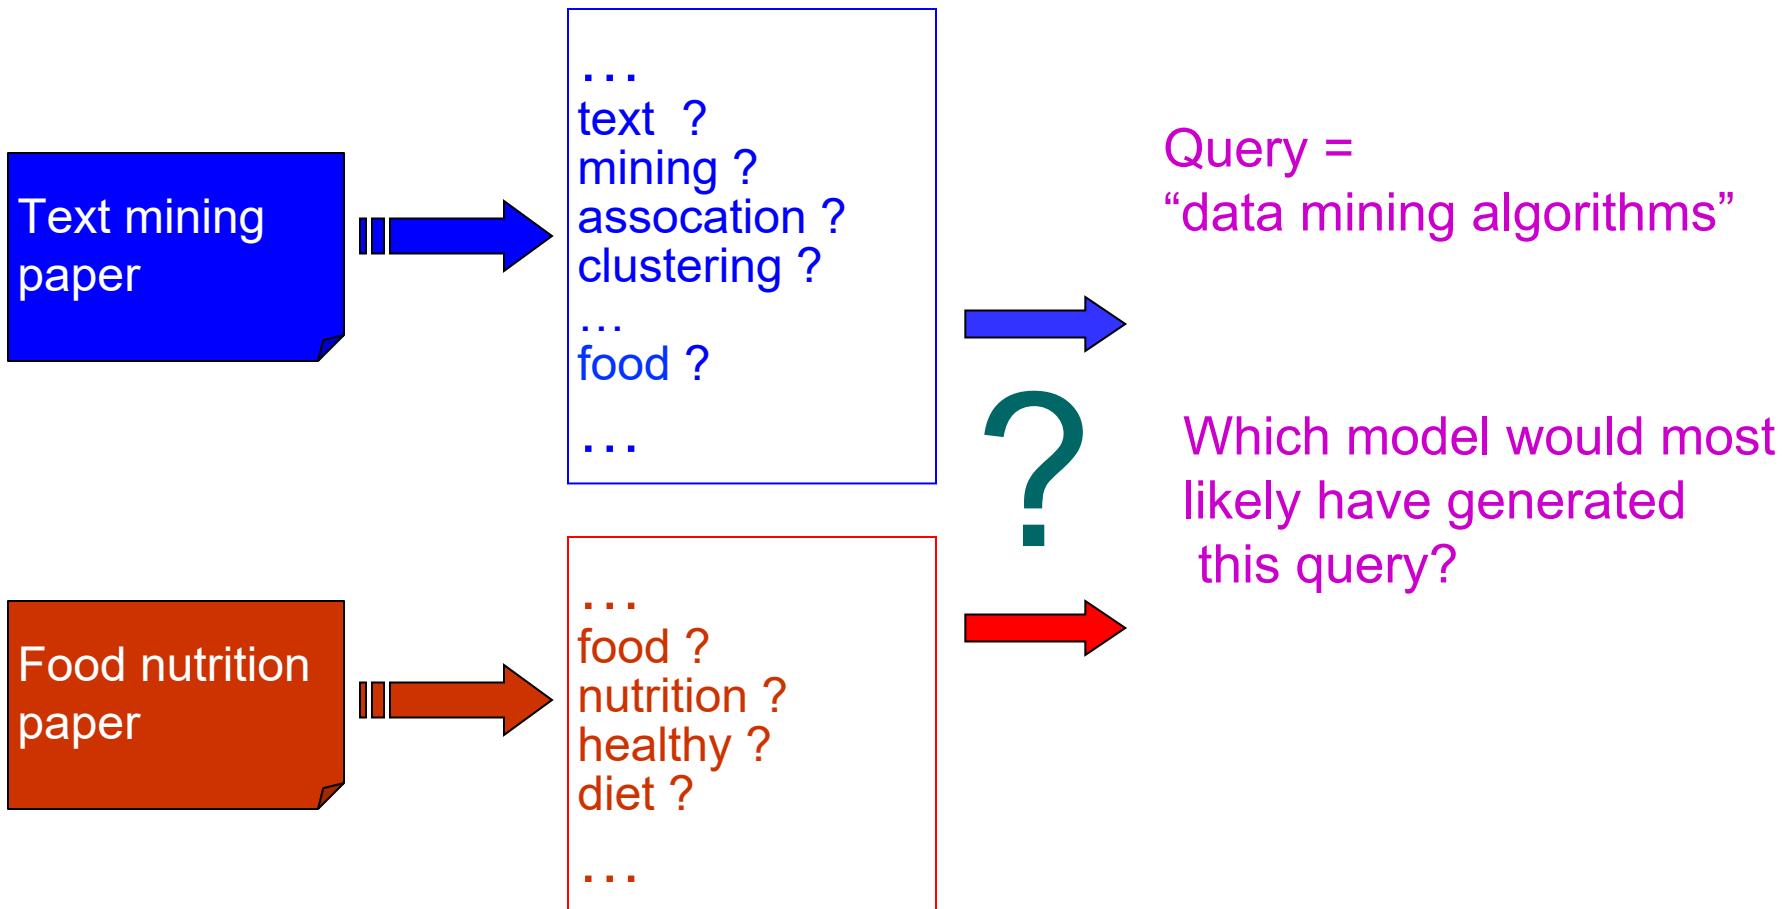
q = Programming text



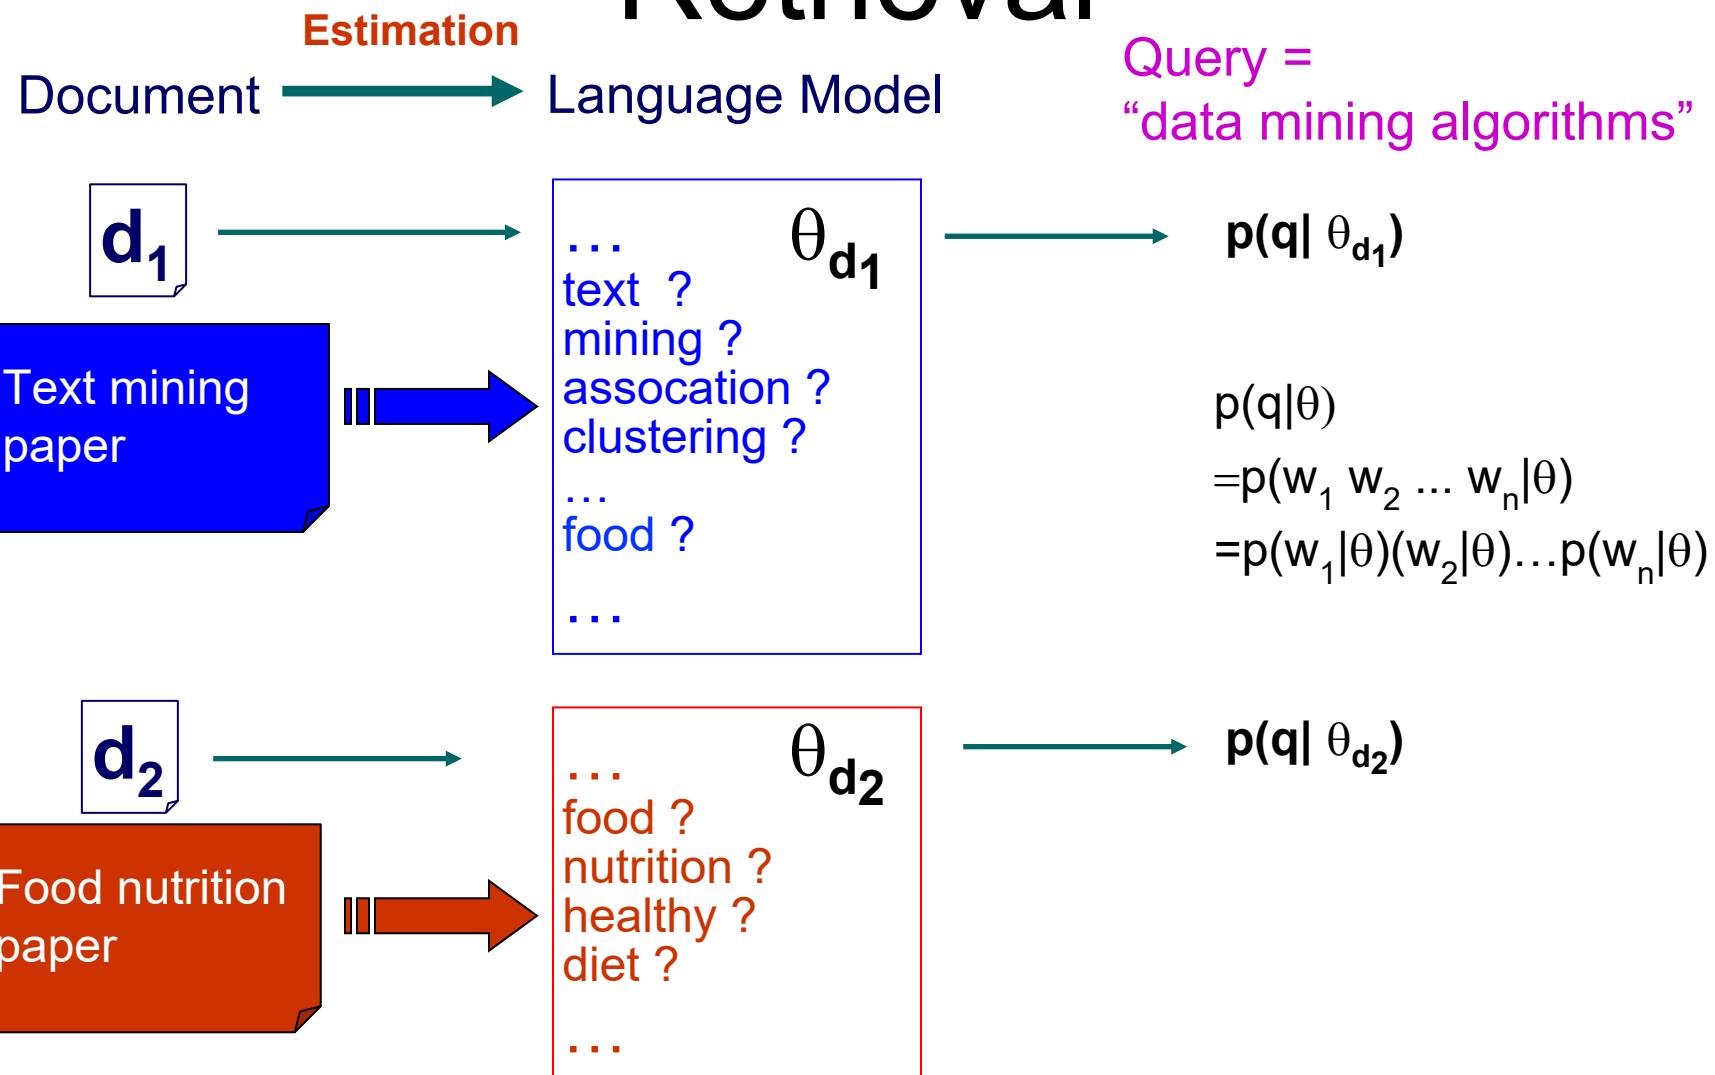
Language Models for IR

Language Models for Retrieval

Estimation
Document → Language Model

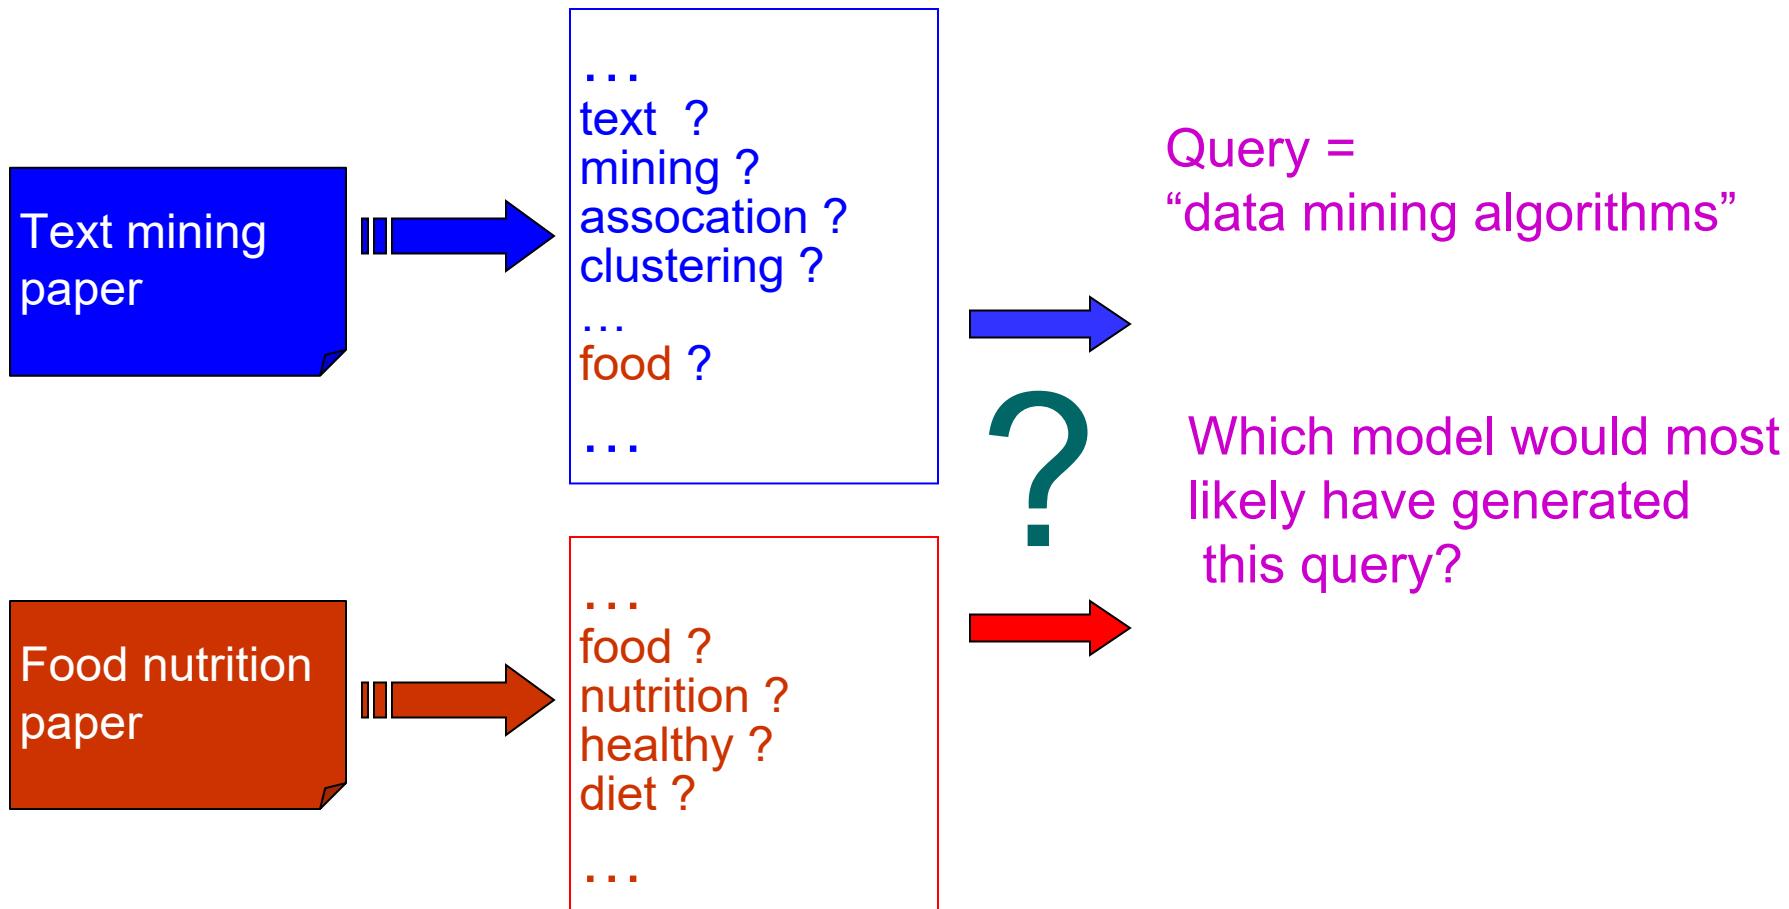


Language Models for Retrieval

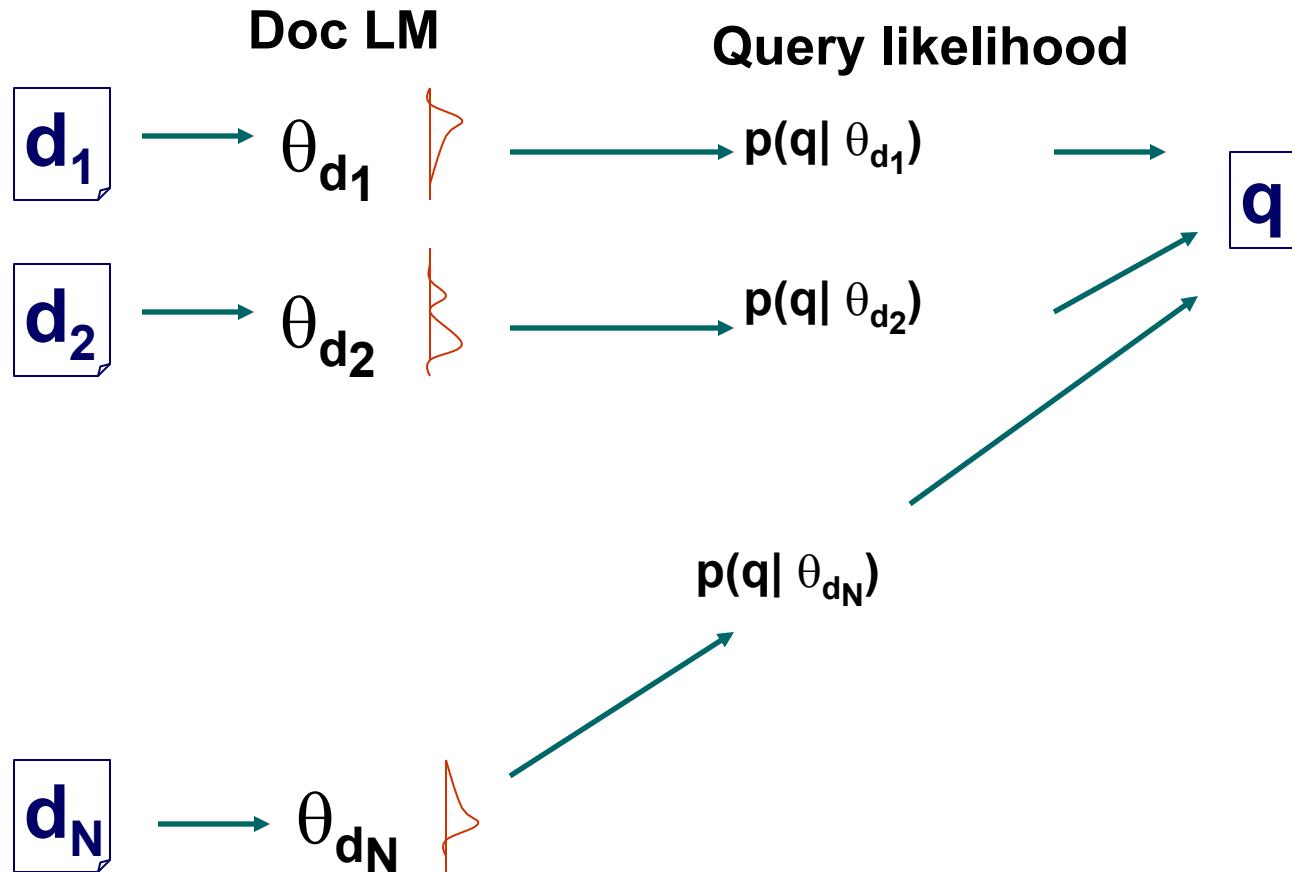


Language Models for Retrieval

Document → Language Model
Estimation



Ranking Docs by Query Likelihood



Retrieval as Language Model Estimation

- Document ranking based on **query likelihood**

$$\log p(q | d) = \sum_i \log p(w_i | d)$$

where, $q = w_1 w_2 \dots w_n$

Document language model

- Retrieval problem \approx Estimation of $p(w_i | d)$

Problem with the Maximum Likelihood Estimator

- What if a query word doesn't appear in a document?
 $p(\text{"information retrieval umiich"} | \theta)$
- In general, what probability should we give a word that has not been observed?
- If we want to assign non-zero probabilities to such words, we'll have to discount the probabilities of observed words
- This is what “smoothing” is about ...

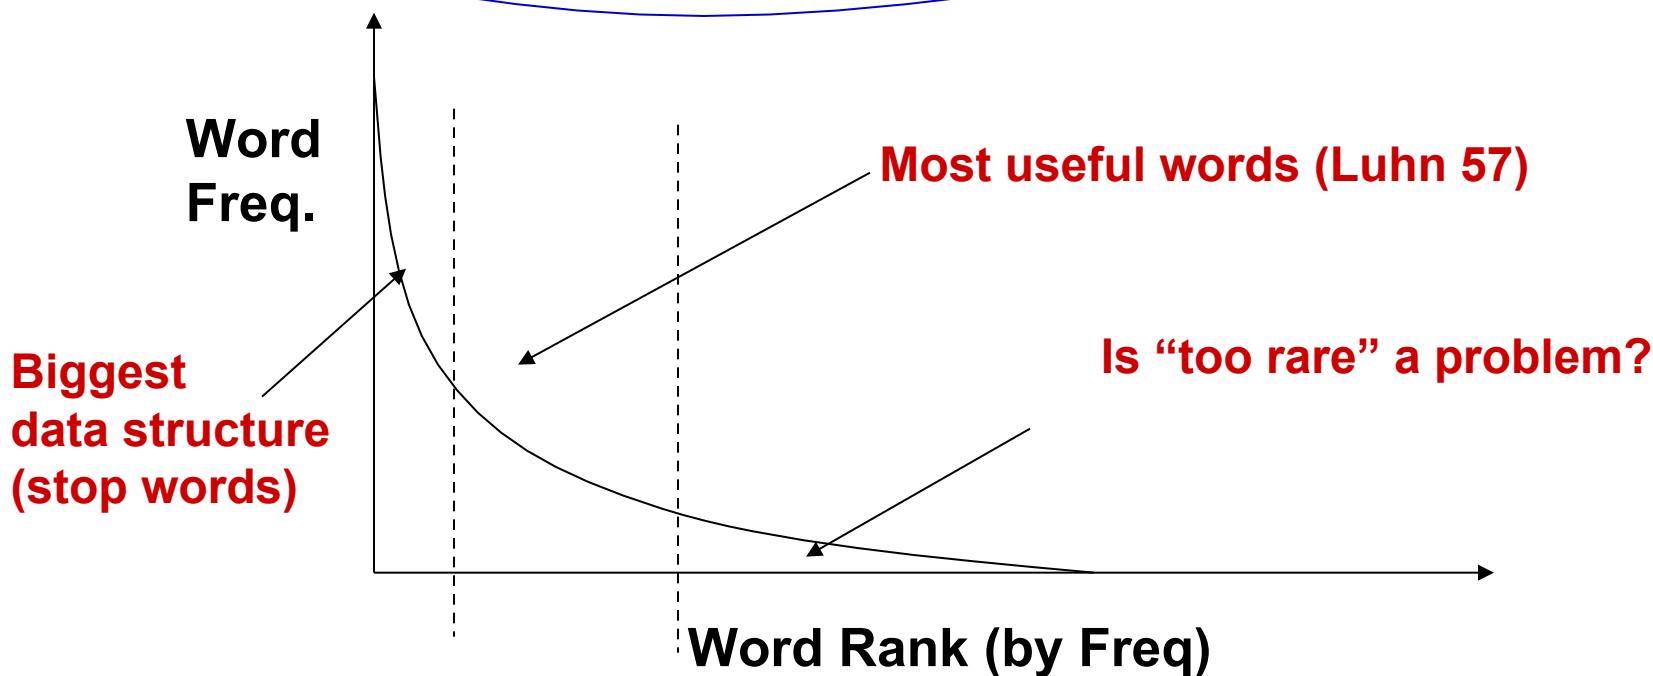
Empirical distribution of words

- There are stable language-independent patterns in how people use natural languages
- A few words occur very frequently; most occur rarely. E.g., in news articles,
 - Top 4 words: 10~15% word occurrences
 - Top 50 words: 35~40% word occurrences
- The most frequent word in one corpus may be rare in another

Revisit: Zipf's Law

- rank * frequency \approx constant

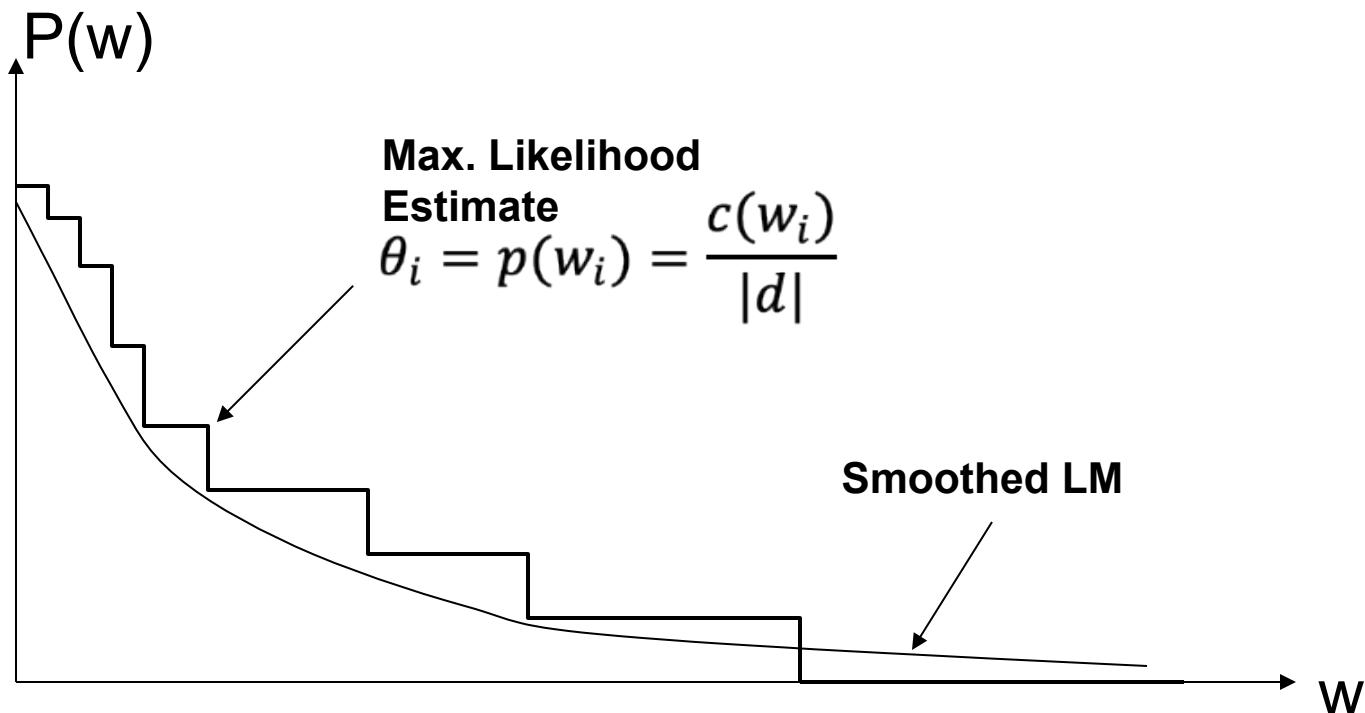
$$F(w) = \frac{C}{r(w)^\alpha} \quad \alpha \approx 1, C \approx 0.1$$



Generalized Zipf's law: $F(w) = \frac{C}{[r(w) + B]^\alpha}$ **Applicable in many domains**

Language Model Smoothing (Illustration)

平滑 (Smoothing) , 它是语言模型里非常重要的一步，主要目的是解决“某个词没在文档里出现导致概率为 0”的问题。



How to Smooth?

- All smoothing methods try to
 - discount the probability of words seen in a document
 - re-allocate the extra counts so that unseen words will have a non-zero count
- Method 1 (Additive smoothing): Add a constant to the counts of each word

平滑方法通过：

- 降低 (discount) 已出现过的词的概率
- 把“省下的”概率分配给没出现的词，让它们概率 > 0

Counts of w in d

$$p(w|d) = \frac{c(w,d) + 1}{|d| + |V|}$$

Length of d (total counts)

“Add one”, Laplace smoothing

Vocabulary size

Additive smoothing

Laplace smoothing:
 $\alpha = 1$

保证即便没出现也有非零概率

$$P(w_i) = \frac{c(w_i) + \alpha}{N + V\alpha}$$

保证概率和为 1

Same idea for
bigram language
models

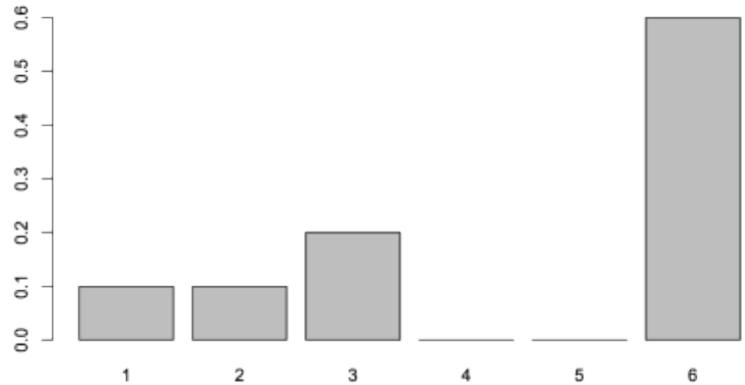
$$P(w_i \mid w_{i-1}) = \frac{c(w_{i-1}, w_i) + \alpha}{c(w_{i-1}) + V\alpha}$$

Smoothing

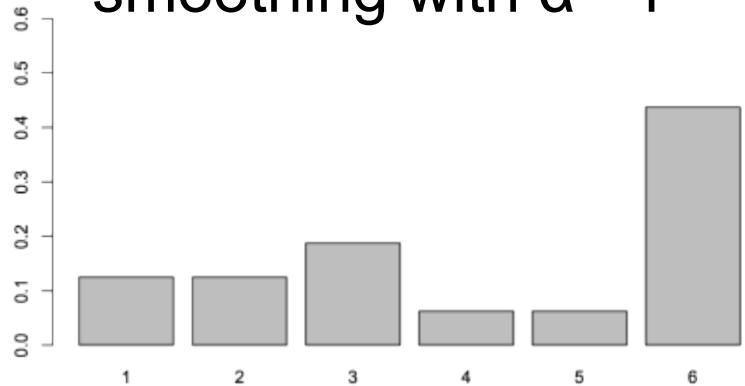
- Maximum likelihood estimates can fail miserably when features are never observed with a particular class.

Smoothing is the re-allocation of probability mass

MLE



smoothing with $\alpha = 1$



How to Smooth? (cont.)

- Should all unseen words get equal probabilities?
- We can use a **reference model** to discriminate unseen words

$$p(w|d) = \begin{cases} p_{seen}(w|d) & \text{if } w \text{ is seen in } d \\ \alpha_d p(w|REF) & \text{otherwise} \end{cases}$$

$$\alpha_d = \frac{1 - \sum_{w \text{ is seen}} p_{seen}(w|d)}{\sum_{w \text{ is unseen}} p(w|REF)}$$

Discounted ML estimate

Reference language model

更智能的平滑方法，不是简单给所有没出现的词分配相同概率，而是借助一个“参考语言模型”(reference model)，让分配给未出现词的概率更合理。

Other Smoothing Methods

- Method 2 (Absolute discounting): Subtract a constant δ from the counts of each word

$$p(w|d) = \frac{\max(c(w;d)-\delta, 0) + \delta |d|_u p(w|REF)}{|d|}$$

$\alpha_d = \delta |d|_u / |d|$

A black arrow points from the text "# uniq words" to the term $|d|_u$ in the formula.

- Method 3 (Linear interpolation, Jelinek-Mercer): “Shrink” uniformly toward $p(w|REF)$

$$p(w|d) = (1 - \lambda) \frac{c(w,d)}{|d|} + \lambda p(w|REF)$$

Two arrows point to the terms in the formula:

- An arrow labeled "ML estimate" points to the term $c(w,d)/|d|$.
- An arrow labeled "parameter" points to the term λ .

$\alpha_d = \lambda$

Jelinek-Mercer Smoothing

Idea: use linear combination of document-specific LM with background LM (corpus, common language):

$$\hat{p}_j(d) = \lambda \frac{\text{freq}(j, d)}{|d|} + (1 - \lambda) \frac{\text{freq}(j, C)}{|C|}$$

could also consider query log as background LM for query

Parameter tuning of λ by cross-validation with held-out data:

- divide set of relevant (d,q) pairs into n partitions
- build LM on the pairs from n-1 partitions
- choose λ to maximize precision (or recall or F1) on nth partition
- iterate with different choice of nth partition and average

Other Smoothing Methods

(cont.)

- Method 4 (**Dirichlet Prior/Bayesian**):
 - Assume adding pseudo counts $\mu(w|REF)$
 - A dynamic coefficient interpolation w.r.t. $|d|$

$$p(w|d) = \frac{c(w;d) + \mu p(w|REF)}{|d| + \mu} = \frac{|d|}{|d| + \mu} \frac{c(w,d)}{|d|} + \frac{\mu}{|d| + \mu} p(w|REF)$$

↑
parameter

How to rank using Dirichlet Smoothing

μ is a hyper parameter that lets us control how much smoothing to do

$$p(w|d) = \frac{c(w,d) + \mu p(w|C)}{|d| + \mu} = \frac{|d|}{|d| + \mu} \frac{c(w,d)}{|d|} + \frac{\mu}{|d| + \mu} p(w|C)$$



$$\log p(q|d)$$



Query term frequency

$$\text{Score}(d, q) = \sum_{w \in q, w \in d} c(w, q) \cdot \log\left(1 + \frac{c(w, d)}{\mu \cdot p(w|C)}\right) + |q| \cdot \log \frac{\mu}{|d| + \mu}$$

Only consider words in the query

Doc. Term frequency

Similar to IDF

Doc. length

So, which smoothing method is the best?

- It depends on the data and the task!
- Many other sophisticated smoothing methods have been proposed...
- Cross validation is generally used to choose the best method and/or set the smoothing parameters...
- For retrieval, Dirichlet prior performs well...

A General Smoothing Scheme

- All smoothing methods try to
 - discount the probability of words seen in a doc
 - re-allocate the extra probability so that unseen words will have a non-zero probability
- Most use a **reference model** (collection language model) to discriminate unseen words

$$p(w | d) = \begin{cases} p_{\text{seen}}(w | d) & \text{if } w \text{ is seen in } d \\ \alpha_d p(w | C) & \text{otherwise} \end{cases}$$

Discounted ML estimate

Collection language model

Comparison of Three Smoothing Methods

(Zhai & Lafferty 01)

- Simplified Jelinek-Mercer: Shrink uniformly toward $p(w|C)$

$$p(w|d) = (1 - \lambda)p_{ml}(w|d) + \lambda(p|C)$$

- Dirichlet prior (Bayesian): Assume pseudo counts $\mu p(w|C)$

$$p(w|d) = \frac{c(w;d) + \mu p(w|C)}{|d| + \mu} = \frac{|d|}{|d| + \mu} p_{ml}(w|d) + \frac{\mu}{|d| + \mu} p(w|C)$$

- Absolute discounting: Subtract a constant δ

$$p(w|d) = \frac{\max(c(w;d) - \delta, 0) + \delta |d|_u p(w|C)}{|d|}$$

① Jelinek-Mercer 平滑 (线性插值)

$$p(w|d) = (1 - \lambda)p_{ml}(w|d) + \lambda p(w|C)$$

- $p_{ml}(w|d)$: 最大似然估计 (文档内频率)。
- $p(w|C)$: 背景语言模型。
- $\lambda \in [0,1]$: 平衡两部分的权重。
- **特点**: 把文档概率直接往全局概率“拉近” (Shrink uniformly), 适合文档很短的情况。

② Dirichlet Prior 平滑 (贝叶斯方法)

$$p(w|d) = \frac{c(w; d) + \mu p(w|C)}{|d| + \mu} = \frac{|d|}{|d| + \mu} p_{ml}(w|d) + \frac{\mu}{|d| + \mu} p(w|C)$$

- $c(w; d)$: 文档中词 w 的出现次数。
- $|d|$: 文档长度。
- μ : 平滑参数, 相当于加入了 μ 个“伪计数 (pseudo counts)”。
- **特点**: 文档越长, 越信任文档自己的统计; 文档越短, 越依赖全局统计。

③ Absolute Discounting 平滑 (绝对折扣)

$$p(w|d) = \frac{\max(c(w; d) - \delta, 0) + \delta |d|_u p(w|C)}{|d|}$$

- δ : 从每个已出现词的计数中减去一个固定值。
- $|d|_u$: 文档中不同词的数量 (unique terms)。
- 被减去的概率总和再分配给未见词, 按背景模型分配。

特点: 相比 Jelinek-Mercer 和 Dirichlet, 这种方法对高频词影响较小, 只是简单减去一个常数。

Evaluation

How do we know if our language model captures
queries or documents?

Evaluation: How good is our model?

- Does our language model prefer good sentences to bad ones?
 - Assign higher probability to “real” or “frequently observed” sentences
 - Than “ungrammatical” or “rarely observed” sentences?
- We train parameters of our model on a **training set**.
- We test the model’s performance on data we haven’t seen.
 - A **test set** is an unseen dataset that is different from our training set, totally unused.
 - An **evaluation metric** tells us how well our model does on the test set.

Training on the test set

- We can't allow test sentences into the training set
- We will assign it an artificially high probability when we set it in the test set
- “Training on the test set”
- **Bad science!**
- And violates the honor code

Evaluation

- The best evaluation metrics are **external** — how does a better language model influence the application you care about?

Extrinsic evaluation (外部评估) of language models

- Best evaluation for comparing models A and B
 - Put each model in a task
 - IR system, query prediction, etc.
 - Run the task, get an accuracy for A and for B
 - How many times people clicked links
 - How many auto-suggests were used
 - Compare accuracy for A and B

给两个模型 (A 和 B) 同样的任务

运行任务，统计每个模型的效果

Intrinsic Evaluation

- A good language model should judge **unseen real language** to have high probability
- Perplexity (困惑度) = inverse probability of test data, averaged by word. 测试数据的逆概率（取平均后开方），越低说明模型越好。
- Not a good idea for IR.
- To be reliable, the test data must be truly unseen, **including knowledge of its vocabulary** (could have unseen words!).
- **核心思想**：直接测语言模型本身，而不是放到应用里。

Experiment design

	training	development	testing
size	80%	10%	10%
purpose	training models	model selection; hyperparameter tuning	evaluation; never look at it until the very end

Summary of Language Model

- Goal: estimate a language model from a sample text
- When using maximum-likelihood estimator
 - count the number of times each word occurs in S, divide by length
- Smoothing to avoid zero frequencies (and probabilities!)
 - discounting methods: add or subtract a constant, redistribute mass
 - better: interpolate with background probability of a word
 - smoothing has a role similar to IDF in classical models
- Smoothing parameters very important
 - Dirichlet works well for short queries (need to tune the parameter)
 - Jelinek-Mercer works well for longer queries (also needs tuning)
 - Lots of other ideas being worked on

You Should Know

- How probabilistic approaches work in retrieval in general.
- Okapi/BM25 is a probabilistic model, although the function looks very similar to a VSM
- The basic idea of ranking docs by query likelihood (“the language modeling approach”)
- How smoothing is connected with TF-IDF weighting and document length normalization

Additional slides

Binary Independence Model

- Traditionally used in conjunction with PRP
- “**Binary**” = **Boolean**: documents are represented as binary incidence vectors of terms (cf. IIR Chapter 1):
 - $\vec{x} = (x_1, \dots, x_n)$
 - $x_i = 1$ iff term i is present in document x .
- “**Independence**”: terms occur in documents independently
- Different documents can be modeled as the same vector

Binary Independence Model

- Queries: binary term incidence vectors
- Given query q ,
 - for each document d need to compute $p(R|q,d)$.
 - replace with computing $p(R|q,x)$ where x is binary term incidence vector representing d .
(d 是二进制向量 x)
 - Interested only in ranking
- Will use odds and Bayes' Rule:

$$O(R|q, \vec{x}) = \frac{p(R=1|q, \vec{x})}{p(R=0|q, \vec{x})} = \frac{\frac{p(R=1|q)p(\vec{x}|R=1, q)}{p(R=0|q)p(\vec{x}|R=0, q)}}{\frac{p(\vec{x}|q)}{p(\vec{x}|q)}}$$

Binary Independence Model

$$O(R|q, \vec{x}) = \frac{p(R=1 | q, \vec{x})}{p(R=0 | q, \vec{x})} = \frac{p(R=1 | q)}{p(R=0 | q)} \cdot \frac{p(\vec{x} | R=1, q)}{p(\vec{x} | R=0, q)}$$

Constant for a
given query

Needs estimation

- Using **Independence Assumption**:

$$\frac{p(\vec{x} | R=1, q)}{p(\vec{x} | R=0, q)} = \prod_{i=1}^n \frac{p(x_i | R=1, q)}{p(x_i | R=0, q)}$$

$$O(R|q, \vec{x}) = O(R|q) \cdot \prod_{i=1}^n \frac{p(x_i | R=1, q)}{p(x_i | R=0, q)}$$

Binary Independence Model

$$O(R|q, \vec{x}) = O(R|q) \cdot \prod_{i=1}^n \frac{p(x_i | R=1, q)}{p(x_i | R=0, q)}$$

- Since x_i is either 0 or 1:

$$O(R|q, \vec{x}) = O(R|q) \cdot \prod_{x_i=1} \frac{p(x_i = 1 | R=1, q)}{p(x_i = 1 | R=0, q)} \cdot \prod_{x_i=0} \frac{p(x_i = 0 | R=1, q)}{p(x_i = 0 | R=0, q)}$$

- Let $p_i = p(x_i = 1 | R=1, q)$; $r_i = p(x_i = 1 | R=0, q)$;
- Assume, for all terms not occurring in the query ($q_i=0$) $p_i = r_i$

$$O(R|q, \vec{x}) = O(R|q) \cdot \prod_{\substack{x_i=1 \\ q_i=1}} \frac{p_i}{r_i} \cdot \prod_{\substack{x_i=0 \\ q_i=1}} \frac{(1-p_i)}{(1-r_i)}$$

	document	relevant (R=1)	not relevant (R=0)
term present	$x_i = 1$	p_i	r_i
term absent	$x_i = 0$	$(1 - p_i)$	$(1 - r_i)$

Binary Independence Model

$$O(R|q, \vec{x}) = O(R|q) \cdot \prod_{\substack{x_i=q_i=1 \\ \text{All matching terms}}} \frac{p_i}{r_i} \cdot \prod_{\substack{x_i=0 \\ q_i=1}} \frac{1-p_i}{1-r_i}$$

Non-matching query terms

$$O(R|q, \vec{x}) = O(R|q) \cdot \prod_{\substack{x_i=1 \\ q_i=1}} \frac{p_i}{r_i} \cdot \prod_{\substack{x_i=1 \\ q_i=1}} \left(\frac{1-r_i}{1-p_i} \cdot \frac{1-p_i}{1-r_i} \right) \prod_{\substack{x_i=0 \\ q_i=1}} \frac{1-p_i}{1-r_i}$$

$$O(R|q, \vec{x}) = O(R|q) \cdot \prod_{\substack{x_i=q_i=1 \\ \text{All matching terms}}} \frac{p_i(1-r_i)}{r_i(1-p_i)} \cdot \prod_{q_i=1} \frac{1-p_i}{1-r_i}$$

All query terms

Binary Independence Model

$$O(R | q, \vec{x}) = O(R | q) \cdot \prod_{x_i=q_i=1} \frac{p_i(1-r_i)}{r_i(1-p_i)} \cdot \prod_{q_i=1} \frac{1-p_i}{1-r_i}$$

如果某个词在相关文档里出现的概率 p_i 远大于在不相关文档里出现的概率 r_i , 那么它就是一个“好词”, 能强烈提升文档得分。

Constant for each query

Only quantity to be estimated for rankings

Retrieval Status Value:

$$RSV = \log \prod_{x_i=q_i=1} \frac{p_i(1-r_i)}{r_i(1-p_i)} = \sum_{x_i=q_i=1} \log \frac{p_i(1-r_i)}{r_i(1-p_i)}$$